

Algorithms

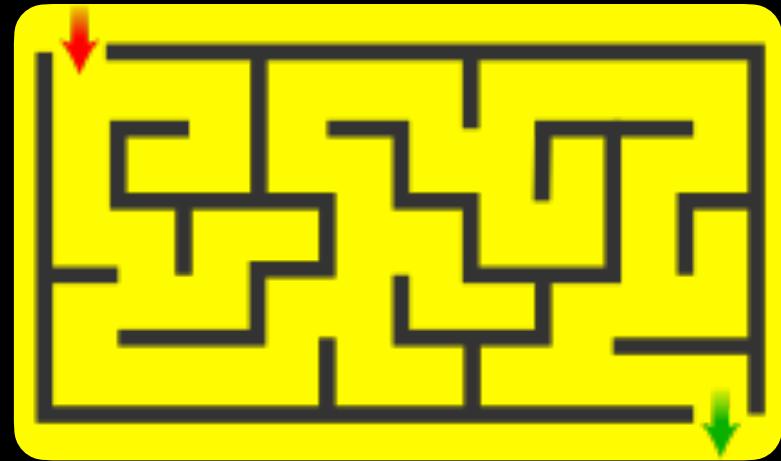
(for Game Design)

Session 4: Navigational Algorithms

Last time

- Mazes
- Maze Creation Algorithms
- Maze Solution Algorithms

Classifications



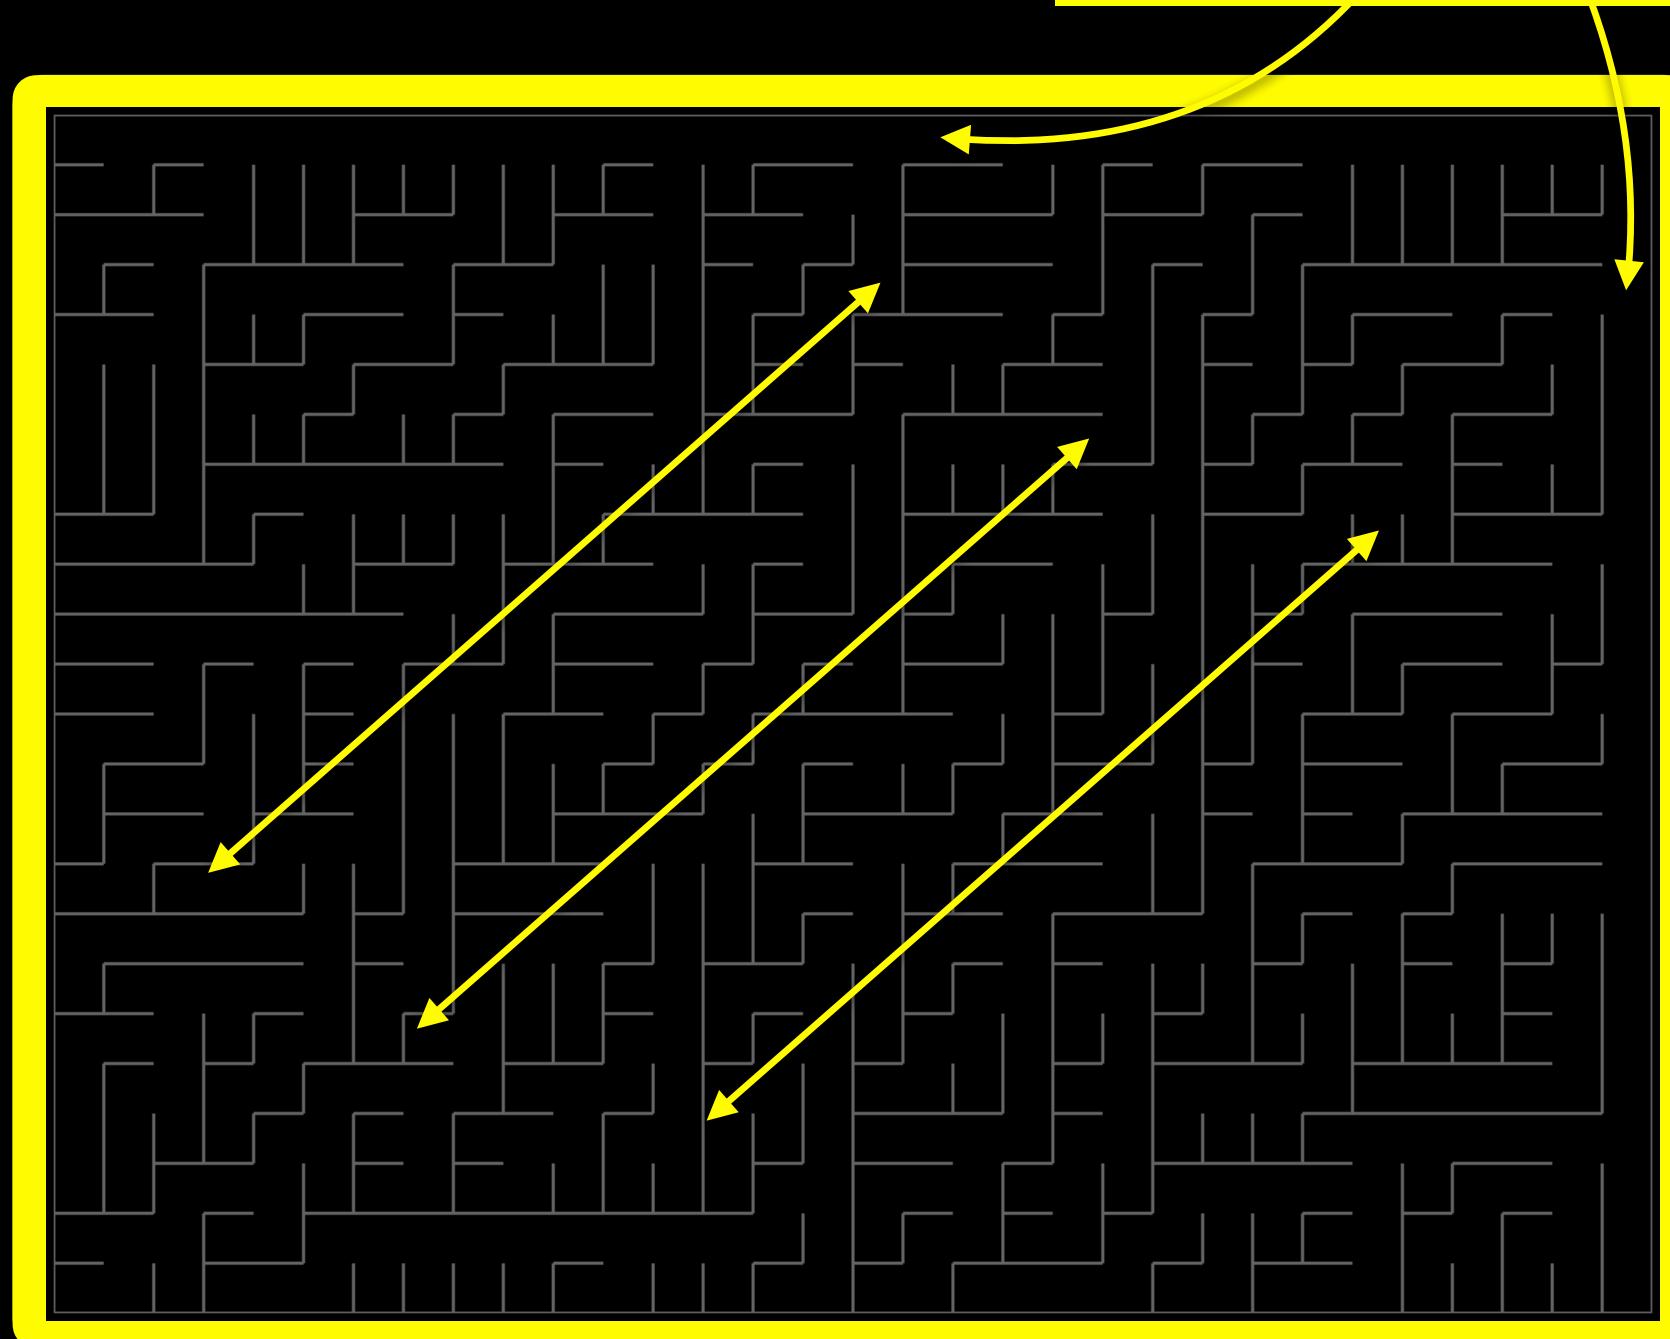
- **Perfect** maze: Has one and only one path between any two cells -- thus plenty of branching, but no loops
 - **Cell**: position in the maze, the smallest geographical element of the maze
- **Unicursal** maze: Has no branches, so there is just a single path from beginning to end



Result: Binary Tree

Top-row and Right-column are always clear

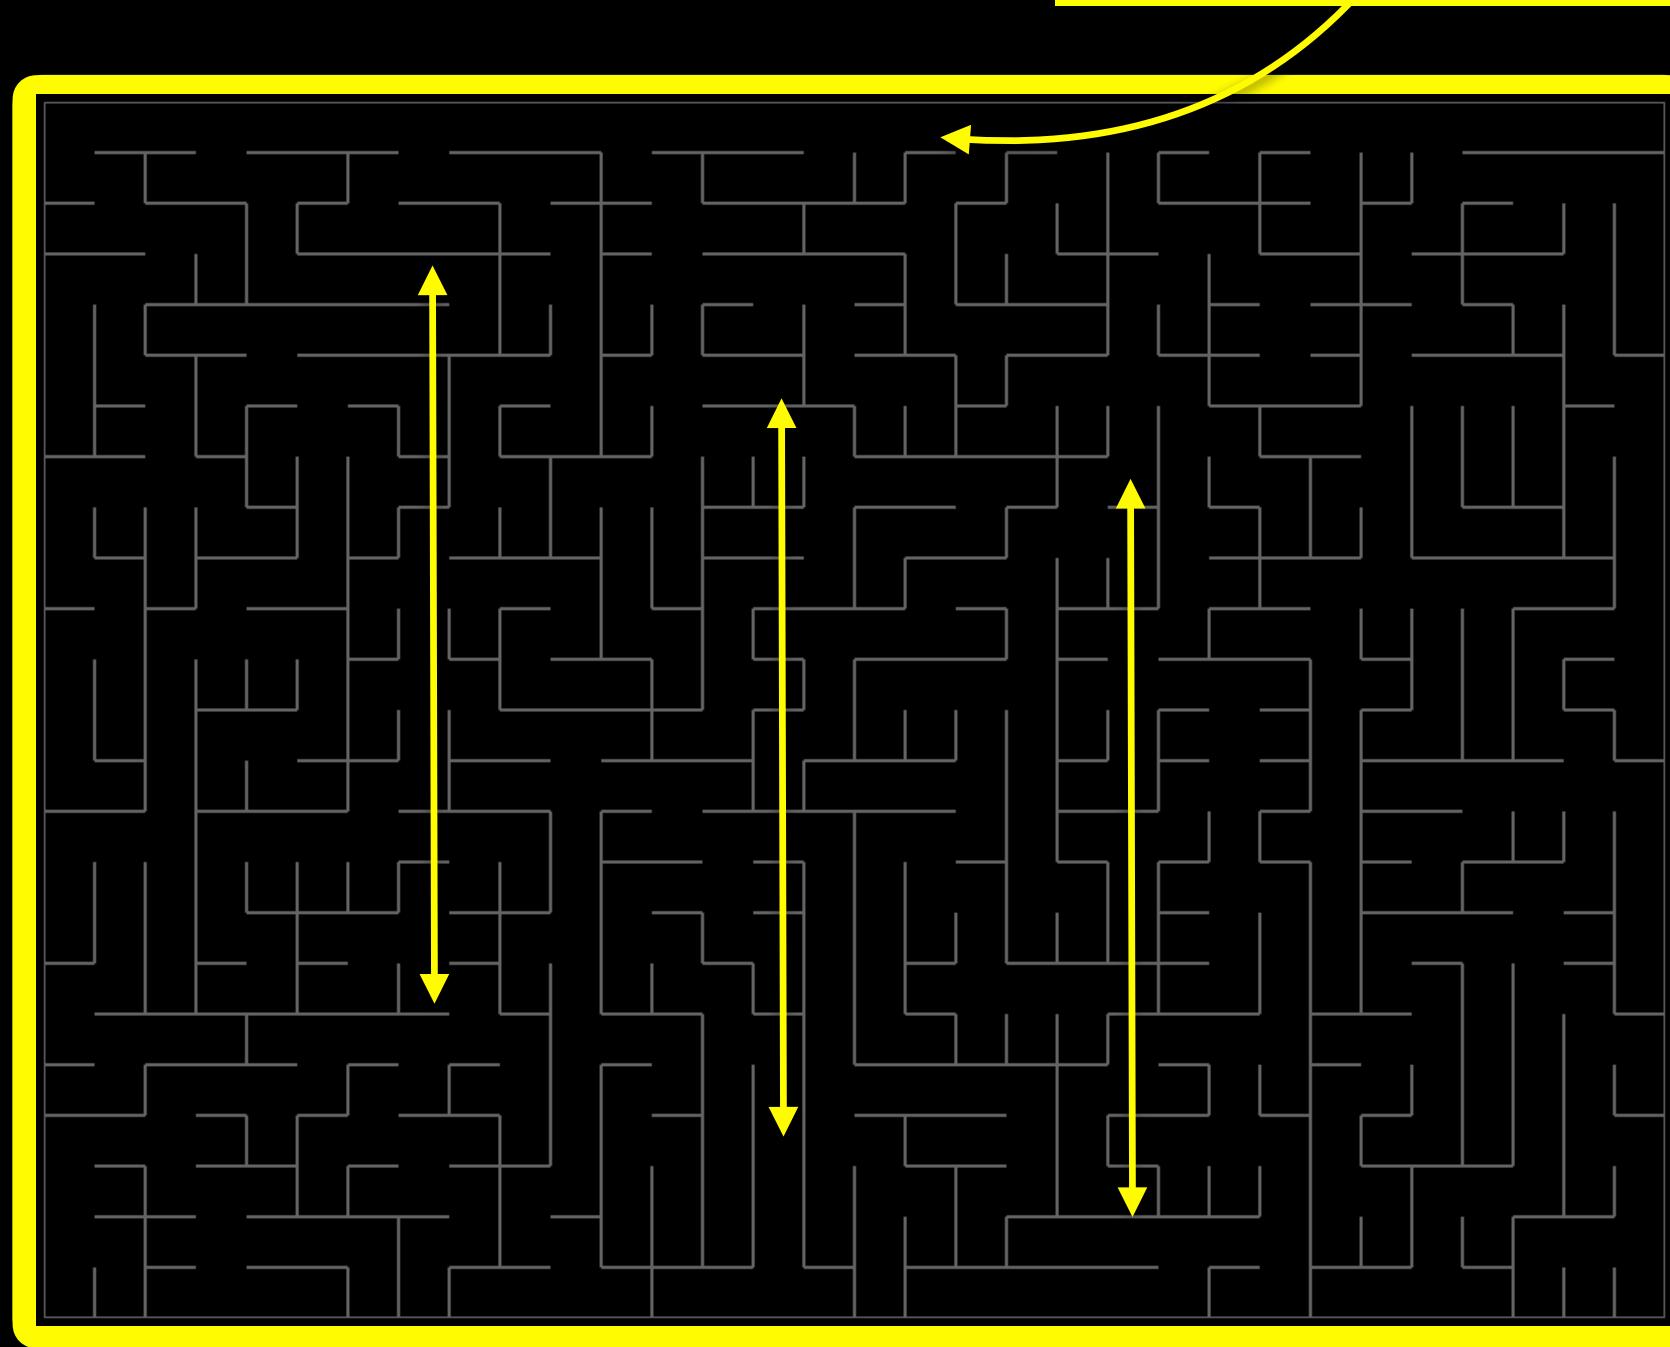
Bias



Result: Sidewinder

Top-row is always clear

Bias



Aldous-Broder

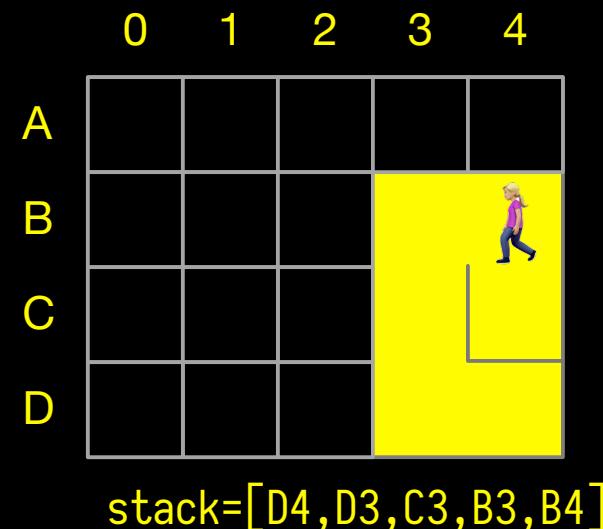
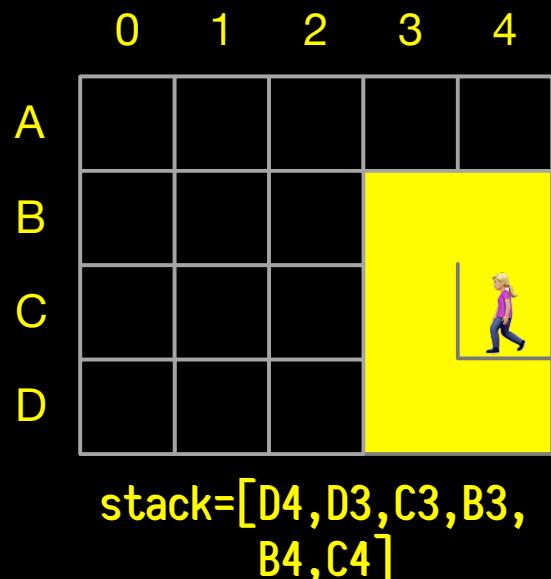
- Discovered simultaneously by two professors with those names
- A **random walk** algorithm: step from cell to random neighbor to create random paths

Wilson's Algorithm

- Also a random walk algorithm
- Also creates mazes with no bias
- Also has inefficient performance characteristics

The Algorithm

- Anytime there is no unvisited neighbor, back up along your path until you do have one, which you pick



Dijkstra's Algorithm

- A very famous graph algorithm, named after a very famous Dutch computer scientist
- Used in network routing algorithms, map navigation, robot movement planning, and many others
- We will start off with a simplified version that works well for our mazes (a "breadth-first search")
 - In sessions to come, we will revisit it several times

Today

- Navigation
- Breadth First Search
- Dijkstra's Algorithm
- Greedy Best First Search
- A*
- Navmeshes

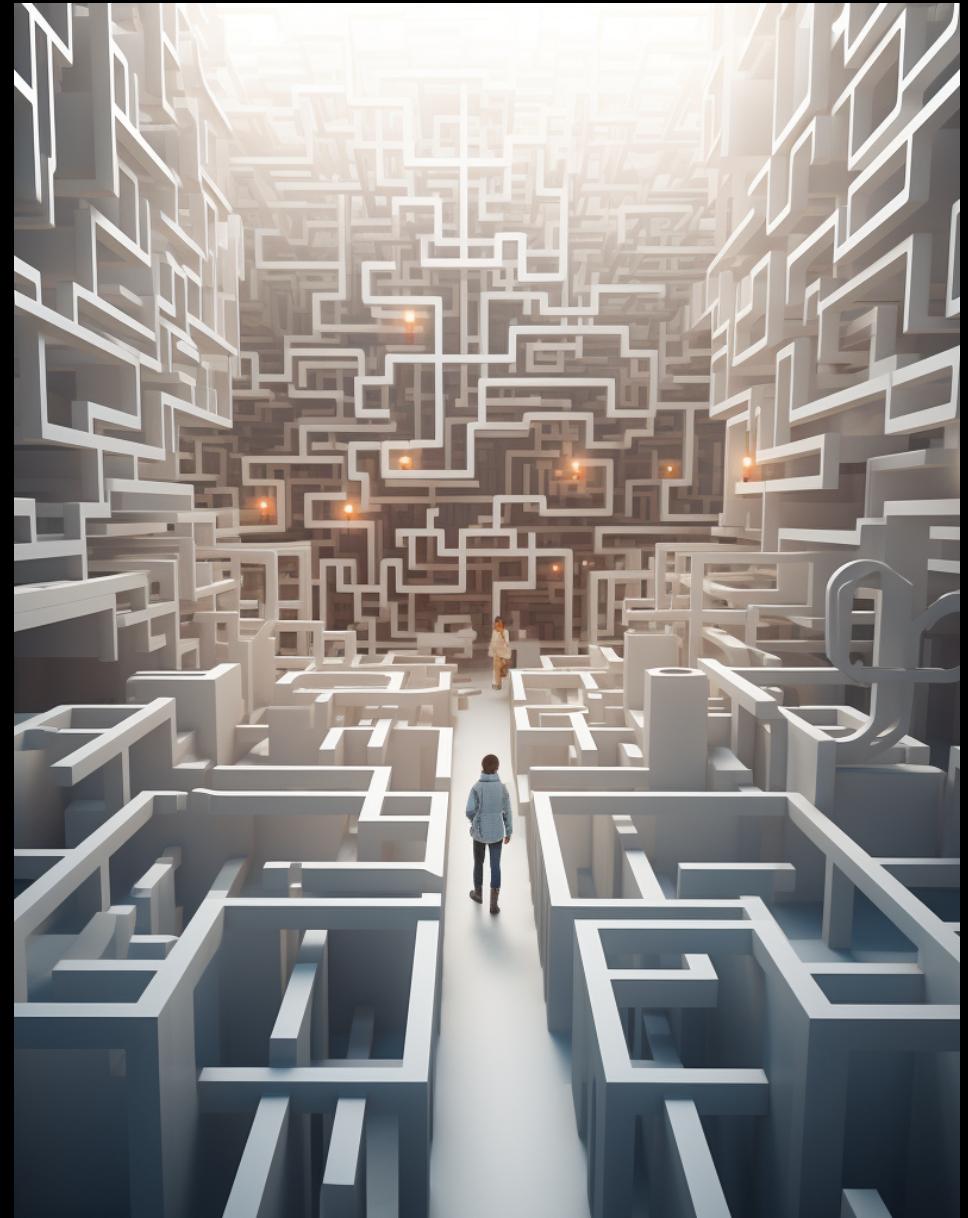
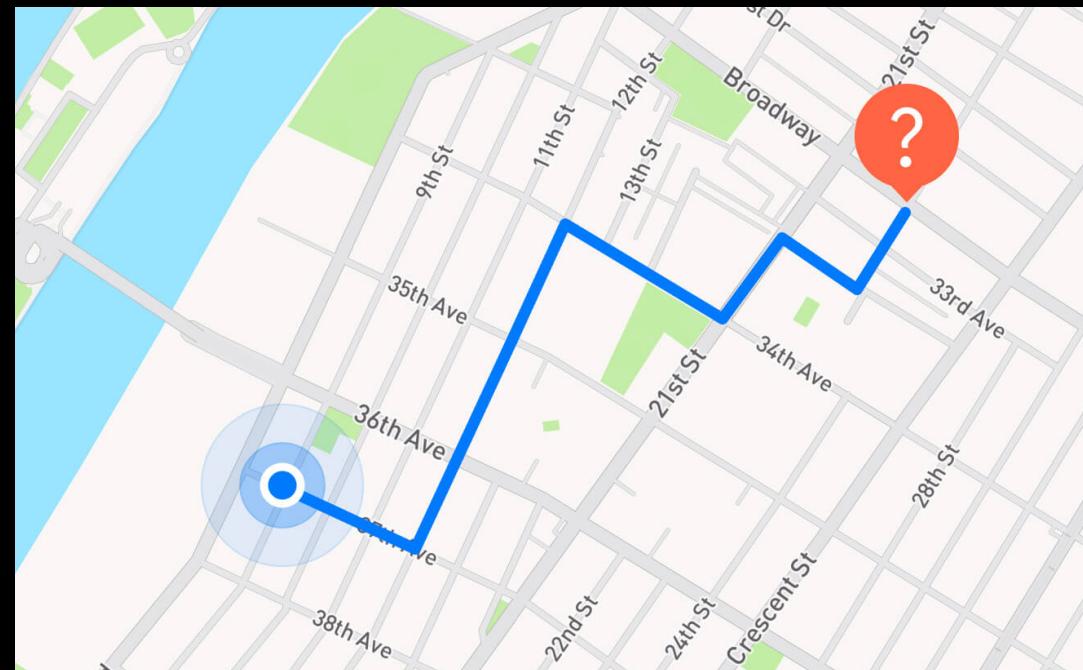


Image credit: Midjourney

navigation

Navigation



- The process of finding a path through the world
 - Often, but not always, about geography

Control

- Often, Navigation is combined with Control
 - Especially in games, where Control means AI
 - Pay attention to "State Machine"
 - Our textbook does this

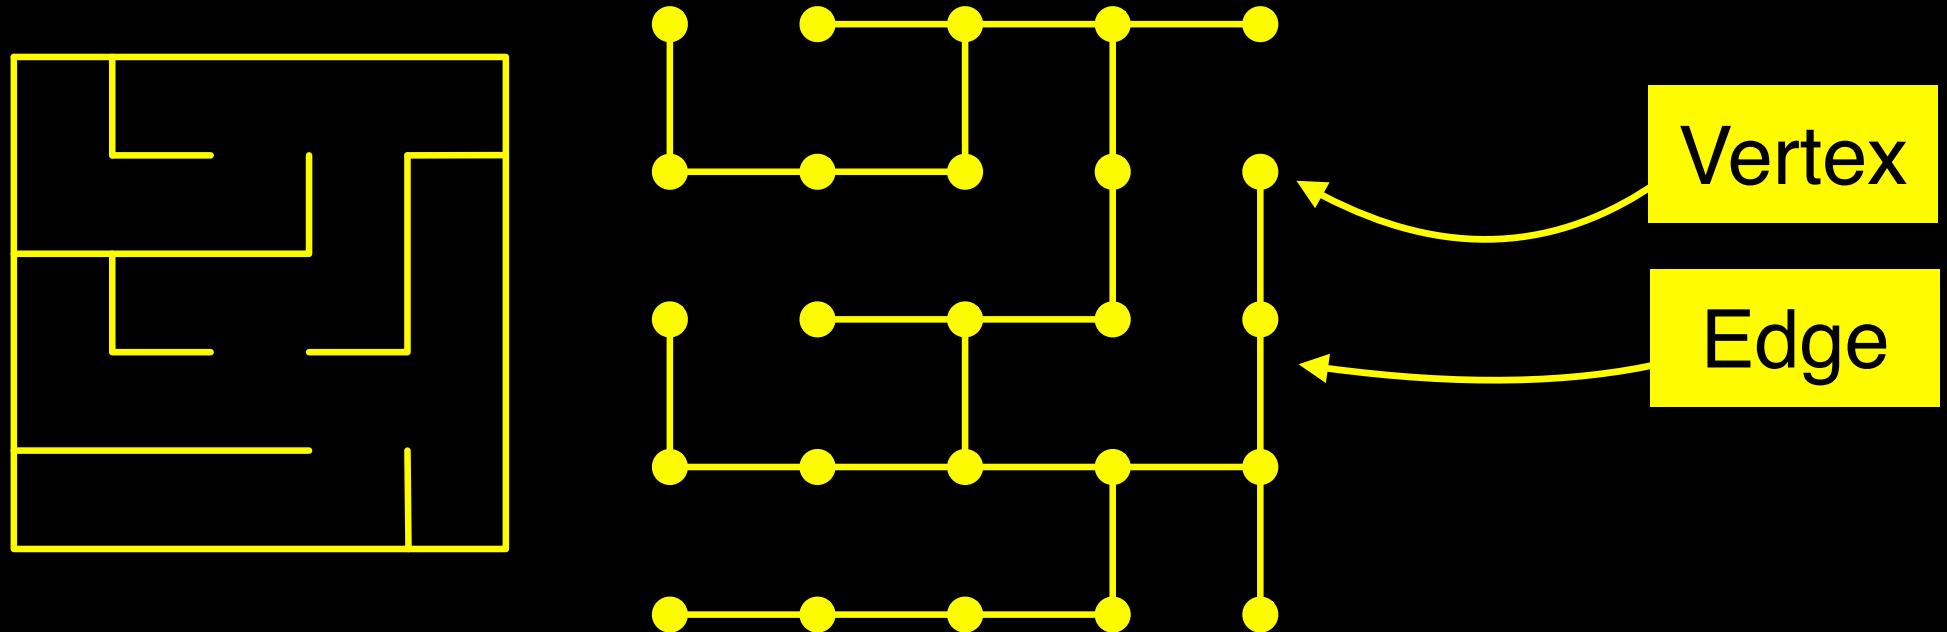
The World

- Consists of locations / ways to get between them
- Mazes → Cells
- Lots of games → Tiles
- In general, a **Graph**

Graph = Set(V) + Set(E)

- **Vertices**: Locations in the world
 - Points, nodes, rooms, cells, towns, planets
 - Annotate Vertex with info about the place
- **Edges**: Connect two Vertices, represent a way to get from one to the other
 - Hallways, tunnels, roads, hyperspace links
 - Annotate Edge with info about the connection (like travel costs)

Mazes are Graphs



- Each cell in a maze → vertex in a graph
- Linked cells → edge in graph

General Graphs

- We will show lots of tiled worlds, mazes today
- But, the algorithms work on any graph

Example: Renowned Explorers

- Uses irregular convex polygons

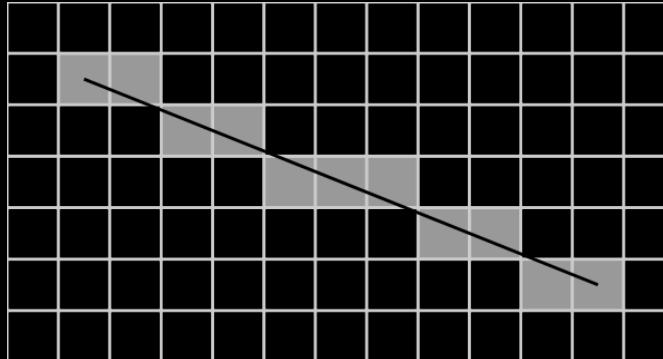


Aside: Game Sizes

- How big should each vertex be?
 - i.e. how much real space should each vertex represent?
- Very interesting blog post about sizes
 - http://www.zarkonnen.com/grids_in_games



Aside: Line Drawing



- **Bresenham's Algorithm**: Classic (1962), low CPU powered (i.e. fast)
- **Linear interpolation** (lerp): Easier to understand
- Orthogonal steps: Modification to either of above such that no diagonal steps are taken
 - No "light leak" through a wall

Back to Navigation

- Navigation: Finding a path from one vertex (V_{start}) to another (V_{end})
- Path: Set of edges (E_0, E_1, \dots, E_n)
 - Connect V_{start} to V_1 to $V_2 \dots$ to V_{end}
 - Generally so $\text{Sum}(E_0.\text{cost}, E_1.\text{cost}, \dots E_n.\text{cost})$ is minimized

Classic Graph Search

- One source, one destination
 - Example: How does my AI monster move towards the player
 - **Greedy Best First Search**
 - **A***

- One source, all destinations (or all source, one destination)
 - Ex: Is it possible to move all my monsters towards the player
 - Ex: Tower defense game, "Flow field"
 - Ex: In which tiles can I place treasure, such that players coming from the entrance can get it
 - **Breadth First Search**: unweighted edges
 - **Dijkstra's Algorithm**: weighted edges, cost > 0
 - Bellman-Ford: weighted edges, cost can be negative

- All sources, all destinations
 - Floyd-Warshall Algorithm
 - Johnson's Algorithm
 - A mixture of Dijkstra's and Bellman-Ford
- Note: All these algorithms (from the past three slides) all work on any general graph
- Simplification: Tiled maps → grid graphs, constant weights, like mazes

breadth first search

Breadth First Search

- We've seen this before
 - Flood fill using a frontier (FIFO queue)
- Continue until frontier is empty
 - Pick the first location from the frontier
 - Mark it as "visited"
 - Add unvisited neighbors to the frontier

```
frontier = [ start ]  
visited = {}  
visited[start] = True
```

```
while not frontier.empty():
```

FIFO. Take from front

```
    current = frontier.pop(0)
```

```
    for neighbor in current.neighbors():
```

```
        if not visited[neighbor]:
```

```
            frontier.append(neighbor)
```

```
            visited[neighbor] = True
```

FIFO. Put on back

Some Pictures

F	56	51	45	39	33	27	20	12	21
57	 A target icon with a red outer ring, a white inner circle, and a black bullseye.	46	40	34	28	22	13	5	14
53	47	41	35	29	23	15	6	1	7
48	42	36	30	24	16	8	2	 A green plus sign symbol composed of four squares forming a cross shape, set against a black background.	4
54	49	43	37	31	25	17	9	3	10
58	55	50	44	38	32	26	18	11	19

Aftermath

- The point of BFS is to discover if you can reach some destination cell
- Test is easy -- just look it up in the visited dictionary

```
def can_reach(some_cell):  
    return some_cell in visited
```

pathfinding with BFS

Pathfinding

- BFS gives a True/False result about whether a destination cell is reachable from the starting cell
- To get the path, need to save a bit more data
 - Usually, we want the shortest path

Modified Algorithm

- For every location we visit, keep track of which neighbor we came from
- I'll change the name of visited to show this

```
frontier = [ start ]  
came_from = {}  
came_from[start] = None
```

Previously, visited

```
while not frontier.empty():  
    current = frontier.pop(0)  
    for neighbor in current.neighbors():  
        if neighbor not in came_from:  
            frontier.append(neighbor)  
            came_from[neighbor] = current
```

Same stuff

Same stuff

Same stuff

This line is new

Pathfinding Results

- Notice that the result is not a path!
 - More like breadcrumbs left on a trail by Hansel and Gretel
 - To use the breadcrumbs, follow the arrows backwards from the target

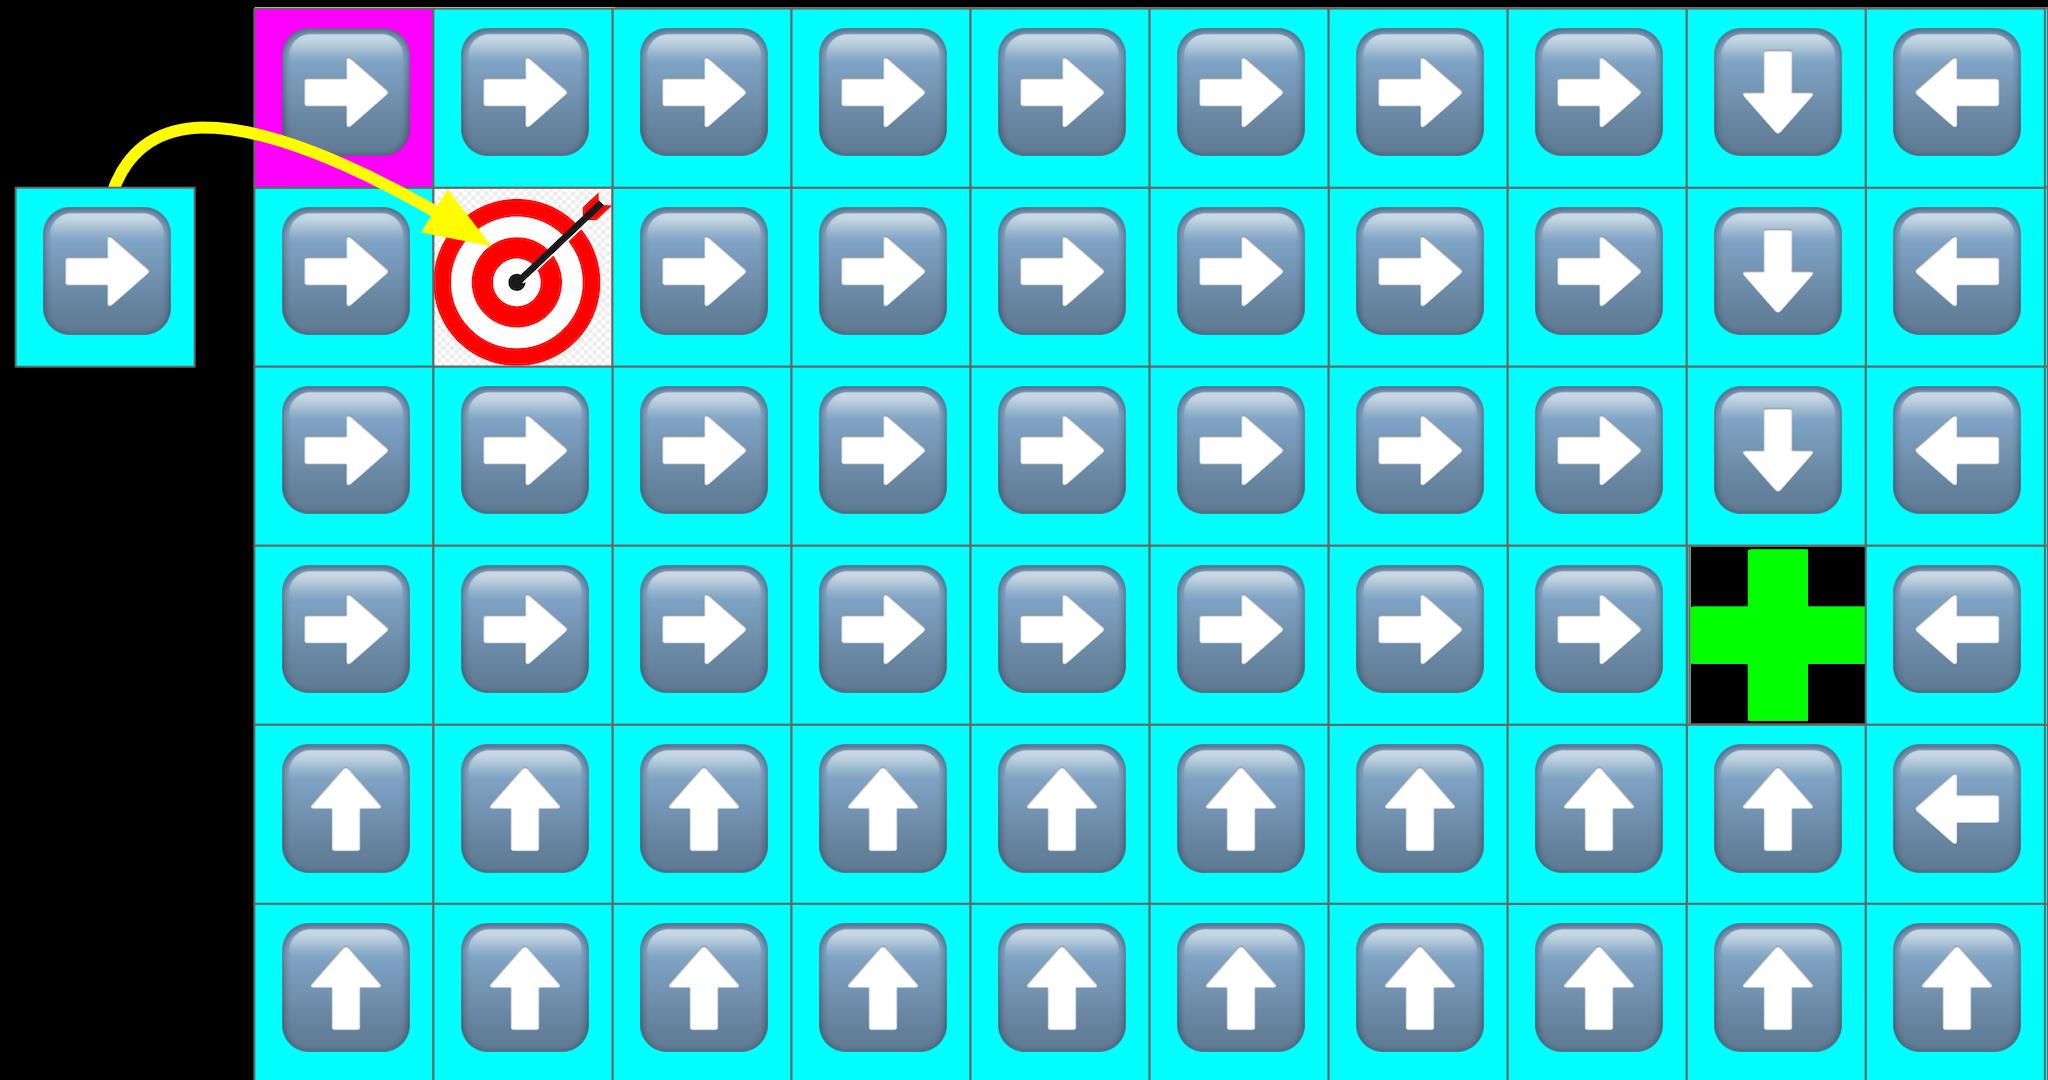
```
def get_path(start, destination, came_from):  
    current = destination  
    path = []  
  
    while current != start:  
        path.insert(0, current)  
        current = came_from[current]  
  
    path.insert(0, start) #optional  
    return path
```

Put the vertex in the front of the list



With this, the list includes start and destination

Pathfinding Pictures



A Few Notes

- Technicality: A path is a sequence of edges in a graph
 - Our path is a sequence of cells (vertices)
 - Equally usable
- The `came_from` dictionary can be used to find a path from ANY cell, not just the target cell
- This is a **flow field**. Very useful for many games, especially tower defense games

early exit

Early Exit

- As shown, algorithm only completes when the entire graph has been explored
- Algorithm can easily be modified to quit when the target is reached

F	56	51	45	39	33	27	20	12	21
57		46	40	34	28	22	13	5	14
53	47	41	35	29	23	15	6	1	7
48	42	36	30	24	16	8	2		4
54	49	43	37	31	25	17	9	3	10
58	55	50	44	38	32	26	18	11	19

		F	45	39	33	27	20	12	21	
			46	40	34	28	22	13	5	14
		F	41	35	29	23	15	6	1	7
	F	42	36	30	24	16	8	2		4
		F	43	37	31	25	17	9	3	10
		F	44	38	32	26	18	11	19	

```
target = some_cell
found_target = False
frontier = [ start ]
came_from = {}
came_from[start] = None

while not frontier.empty():
    current = frontier.pop(0)
    if current == target:
        found_target = True
        break

    for neighbor in current.neighbors():
        if neighbor not in came_from:
            frontier.append(neighbor)
            came_from[neighbor] = current
```

This lets you know, after the end of this code, if the while loop stopped because the target was found or not

obstacles

Obstacles

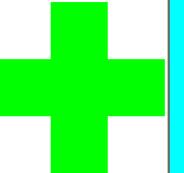


- Typically, you won't be navigating across big caves with nothing in the way
- Mazes had lots of obstacles, walls between unlinked cells
- Navigation algorithms can easily work with obstacles
 - Else, they wouldn't be very useful

Representation

- Goal: Ensure obstacles never get put into the frontier
- Pure graph: obstacles are represented by a lack of edge between two vertices
 - No edge \rightarrow can't get there from here
- Tiles: Don't list obstacles as neighbors
 - Or, have a list of obstacle cells

Pictures

F	47	44		30	26	21	15	22	27
48		41		28	23	16	10		
46	42	38		24	17	11	5	1	6
43	39	35		18	12	7	2		4
40	36	33		19	13			3	8
37	34	32	31	29	25	20		9	14

```
obstacles = [ ] #list of cells that are walls/stuff  
frontier = [ start ]  
came_from = {}  
came_from[start] = None
```

Very easy extension to
our BFS code

```
while not frontier.empty():  
    current = frontier.pop(0)  
    for neighbor in current.neighbors():  
        if neighbor not in came_from and  
            neighbor not in obstacles:  
            frontier.append(neighbor)  
            came_from[neighbor] = current
```

Check for obstacle first

Dijkstra's

Dijkstra's Algorithm

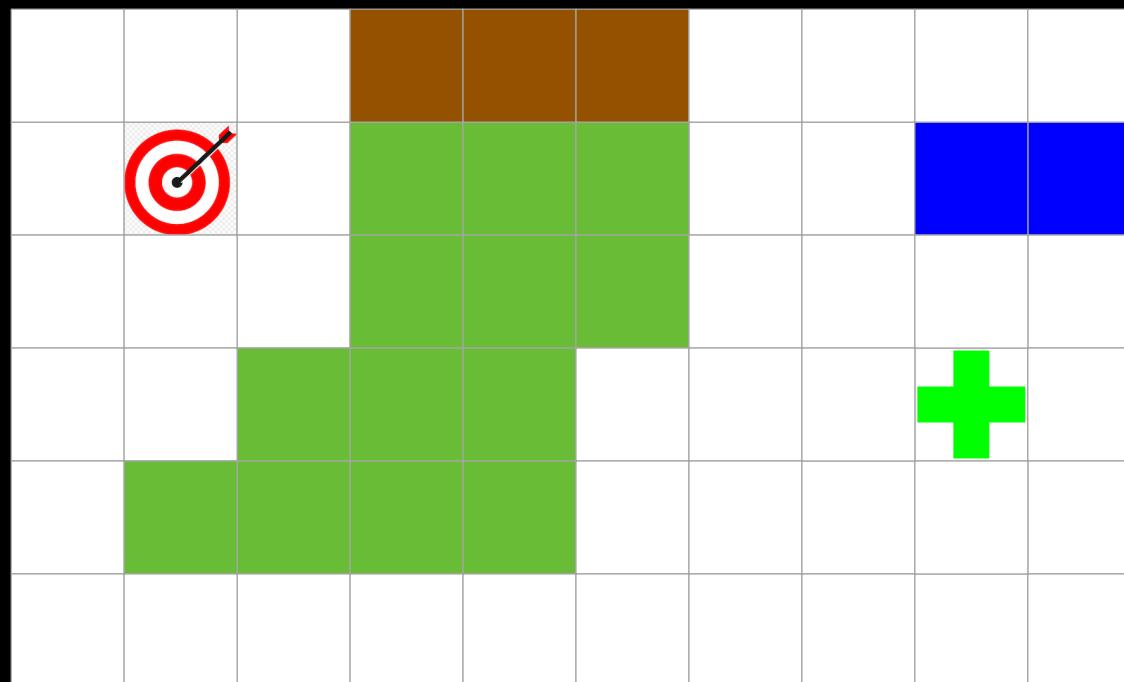
- Calculates the minimum cost path from one vertex to all other vertices
- Implies that edges have costs
 - Terrain costs
 - Does it cost more to go through a door? Or to cross a river?

- Dijkstra's is a "single source, all destinations" algorithm
- The idea:
 - Mark each vertex with the lowest cost to get there
 - Put neighbors into the frontier with additional cost of the edge
 - There may be multiple paths to a particular vertex (which means it will be in the frontier multiple times)
 - Take out the lowest-cost vertex from the vertex
 - This means there can never be a lower cost path discovered to that vertex in the future

Pictures

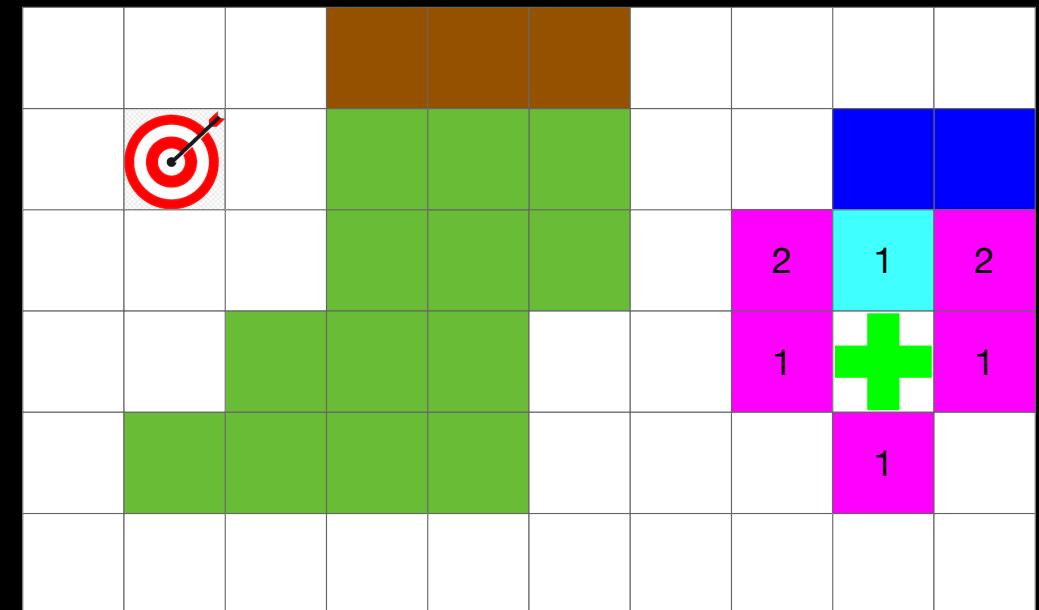
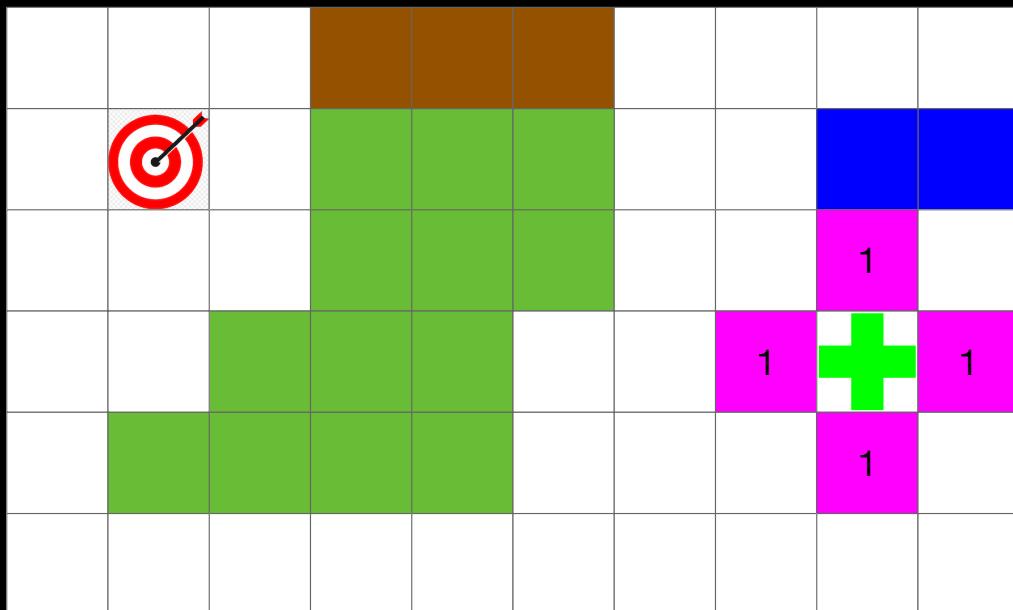
5	1
3	∞

- Let's imagine that our world has less obstacles, but also some forests (movement cost of 3)
- And some mountains (movement cost of 5)



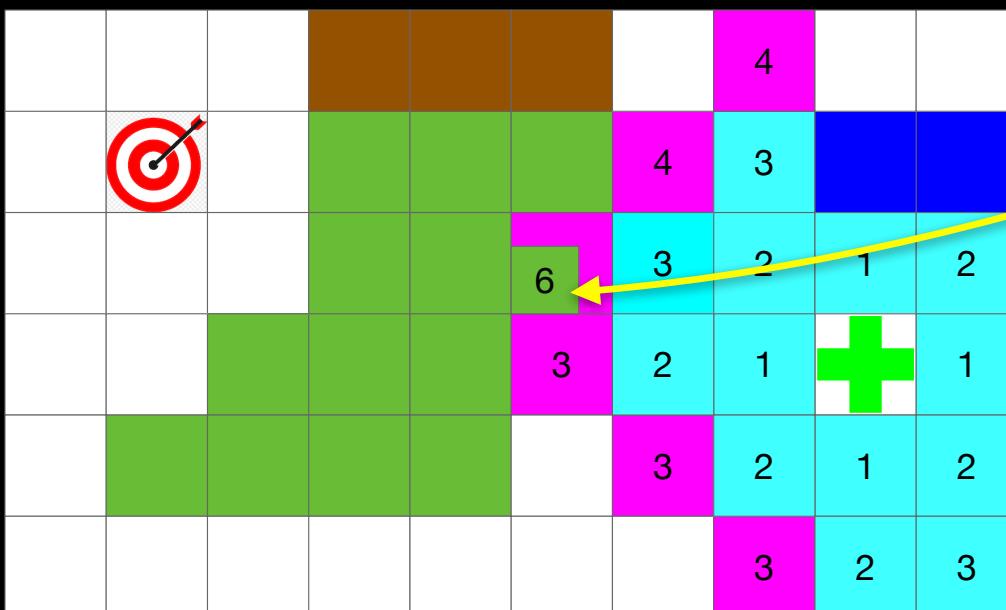
More Pictures

- In the beginning, it looks exactly like we've seen before
- But! We are adding one because that is the terrain cost, not because we always add one



Even More Pictures

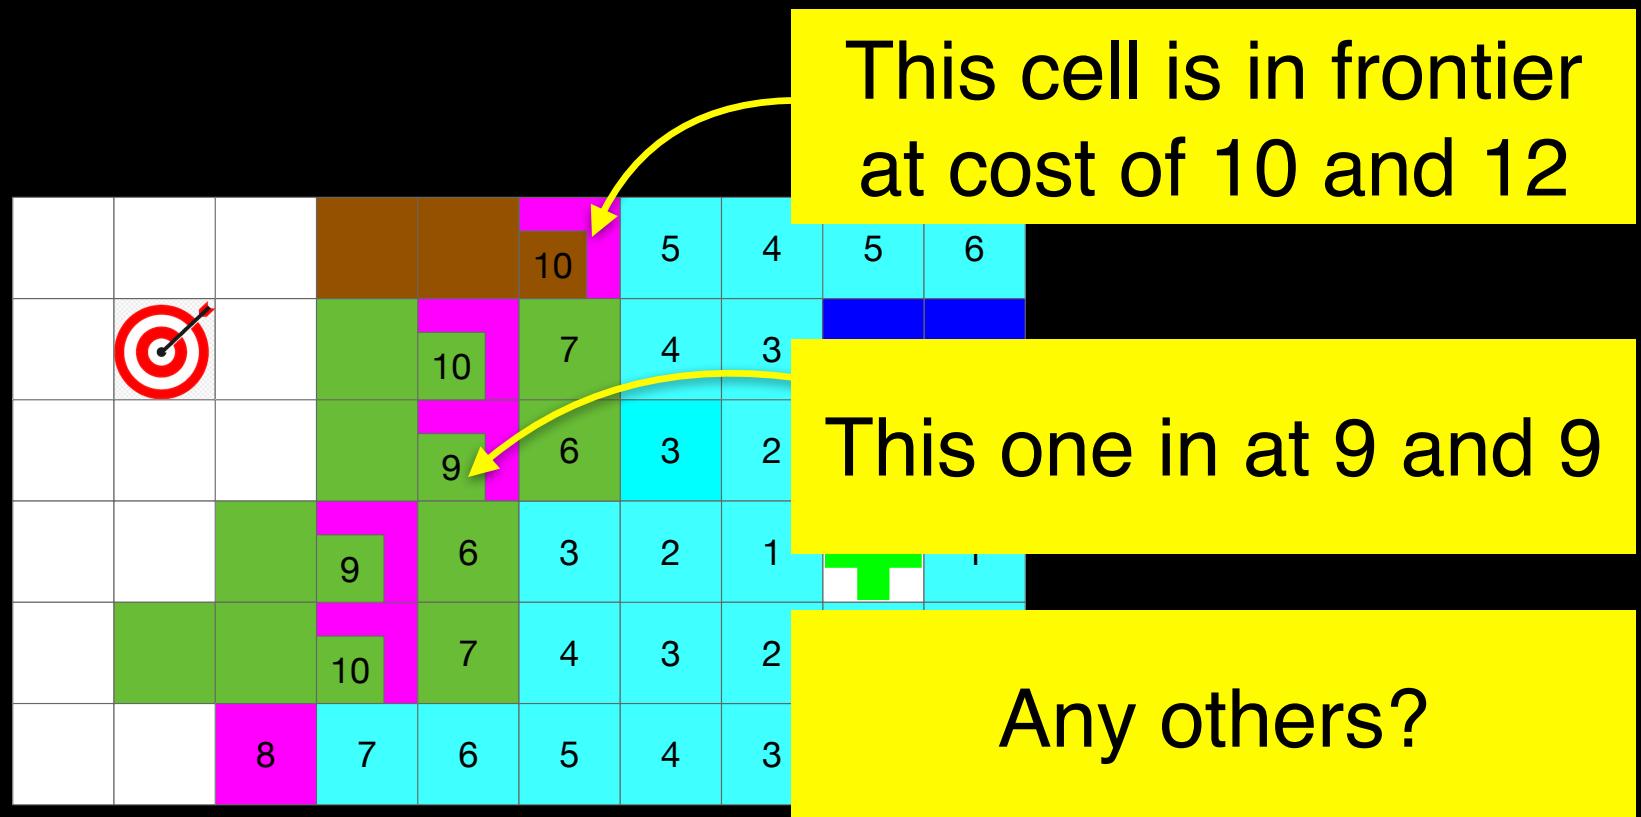
- When the frontier extends to the forest, we add three to the current vertex cost, not one



Notice! Cost of 6, not 4

... Pictures

- Some locations get put into the frontier multiple times, as multiple paths to that location are discovered



Last Pictures

- The path around the forest is pretty close to the cost of the path through the forest
- Here's the final situation -- least cost is 15



Implementation

- The algorithm relies on easily finding the lowest cost of all the cells in the frontier
- Efficiently: use a **Priority Queue** (also called a min-heap)
- In Python, the **heapq** module is part of the standard library
 - Uses a list, but managed as a priority queue

```
import heapq
frontier = [ ]
heapq.heappush(frontier, (0, start, None))
    #tuple of distance, cell, came from cell
came_from = {start : None} #Completed knowledge
```

Tuples kept in sorted order,
so distance must be first

```
while not frontier.empty():
    cost, current, from = heapq.heappop(frontier)
    if current in came_from:
        continue
    came_from[current] = from
    #early exit code could be here

    for neighbor in current.neighbors():
        if neighbor not in came_from:
            new_cost = cost + neighbor.cost
            heapq.heappush(frontier,
                           (new_cost, neighbor, current))
```

Multiple records are likely, so
ignore longer routes

greedy best first search

Efficiency!

- BFS and Dijkstra's spread out to search in all directions
 - If there is only a single goal, this is inefficient
 - Just head towards the goal
- But, if we knew which step was the best way to the goal, we'd already know the entire path

Heuristic

- Can we use a heuristic function to guide our search?
- Heuristic is a "rule of thumb" or a loosely defined guideline → imperfect

```
def heuristic(ax, ay, bx, by):  
    # Manhattan distance on a square grid  
    return abs(ax - bx) + abs(ay - by)
```

- Manhattan distance make sense for a grid of squares → other situations, it might not

Heuristic Search?

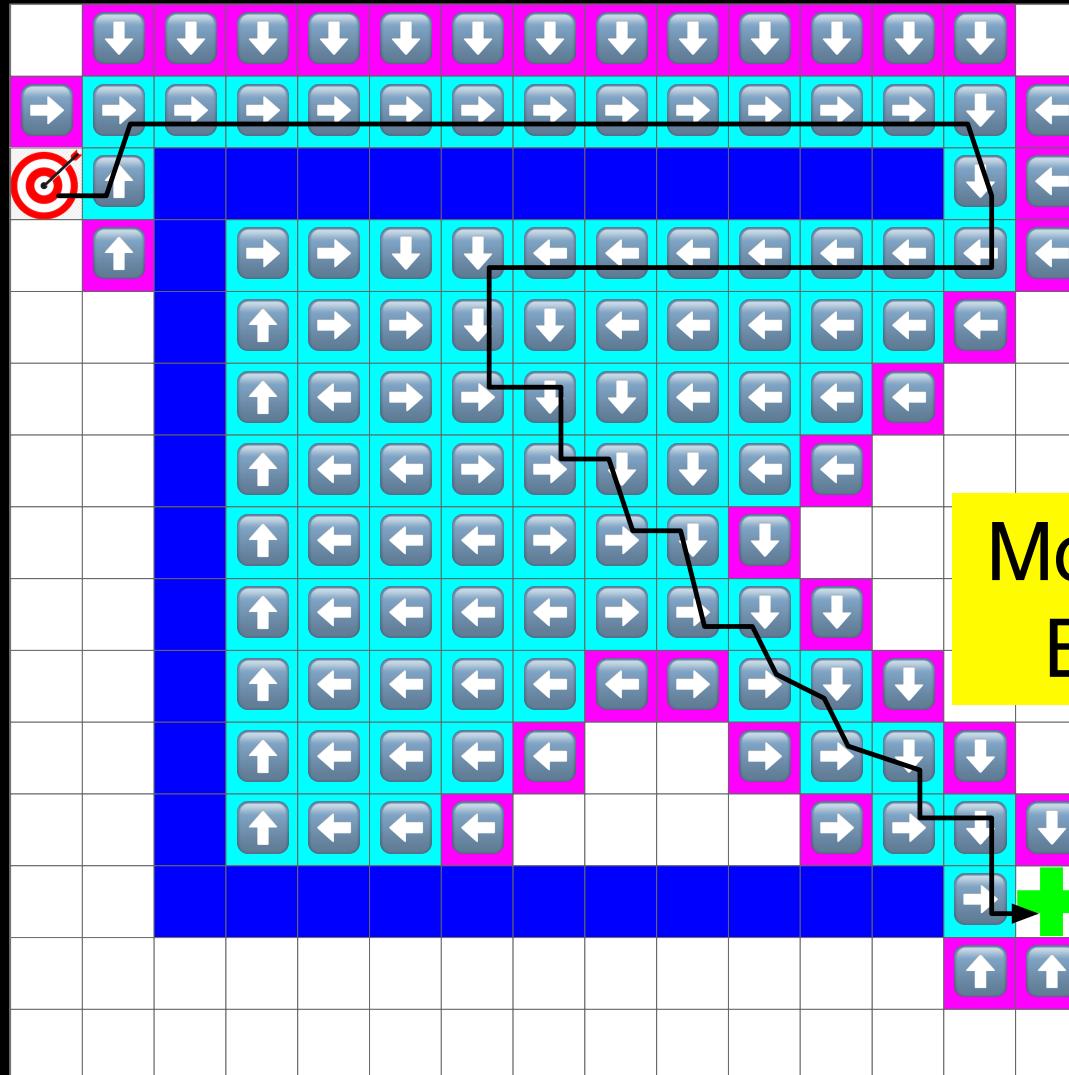
- To use a heuristic is fairly simple
- Keep the frontier / loop structure
- Each time around the loop, take the cell from the frontier with the lowest heuristic distance
 - i.e. the one we think is closest to the goal

```
import heapq
frontier = [ ] #tuple of h_dist, cell, came_from
heapq.heappush(frontier, (0, start, None))
came_from = {}
```

Using the heuristic value just for sorting order

Sounds Fantastic!

Yeah, but ...



More efficient than
BFS / Dijkstra's

Resulting path
is worse than
BFS / Dijkstra's

A^{*}

So Far

- Dijkstra's Algorithm
 - Resulting path is perfect
 - Inefficient -- searches everywhere
 - Greedy Best First Search
 - Efficient
 - Resulting path is not good
- Would a combination be better?
- Tradeoffs?

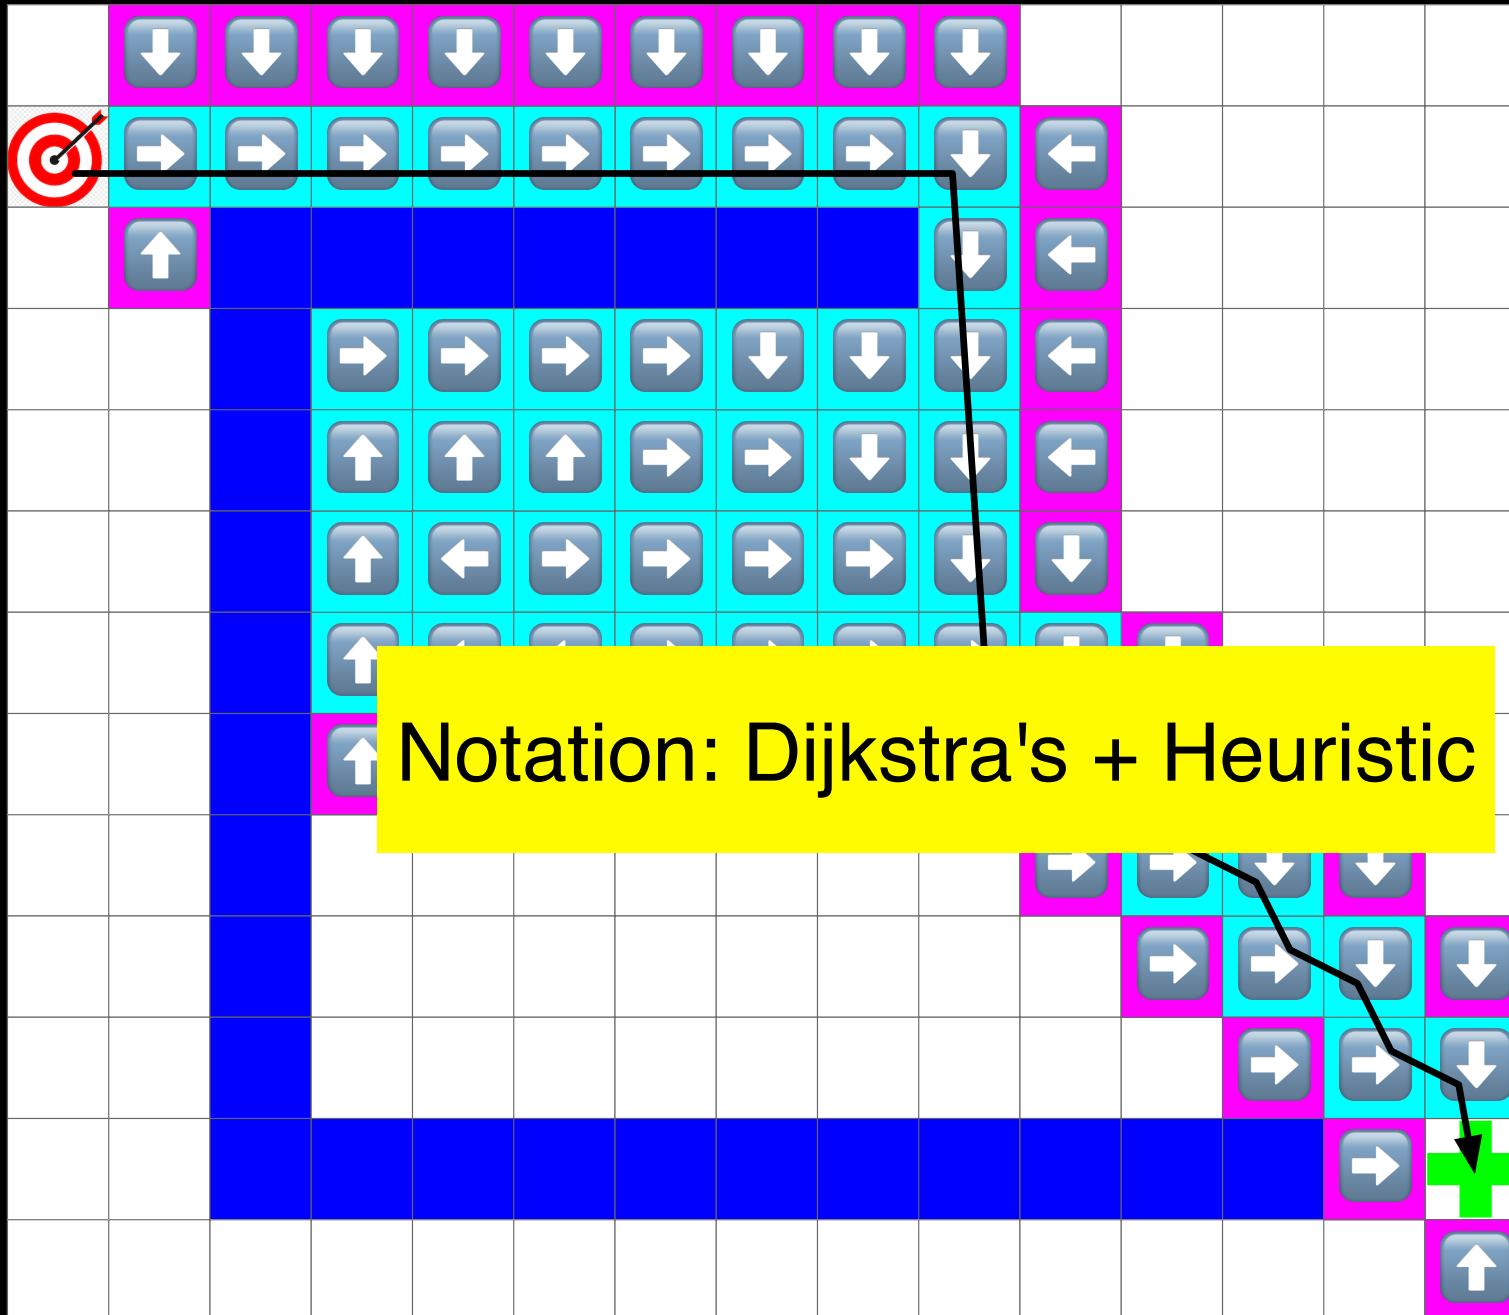
A* Algorithm

- A graph search algorithm
 - Developed at SRI in 1962
 - Developed for robotics
- A best-of-both worlds algorithm, combining Dijkstra's and GBFS

Some comments

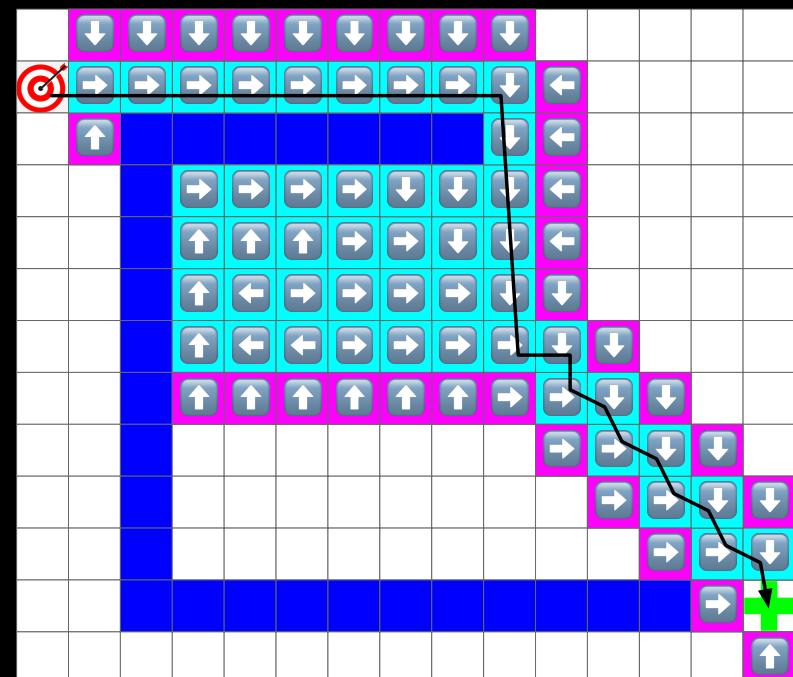
- If GBFS would have found the right path, A* will explore exactly the same area
- But, if GBFS → longer path, then A* is restrained by Dijkstra distance and keeps to the right path
- Dijkstra's algorithm is simply A* with a constant heuristic value
- A* uses the heuristic to re-order the search so it finds the goal sooner

Pictures



A* Results

- Search area doesn't look much different from GBFS, does it?
- But, resulting path is much better



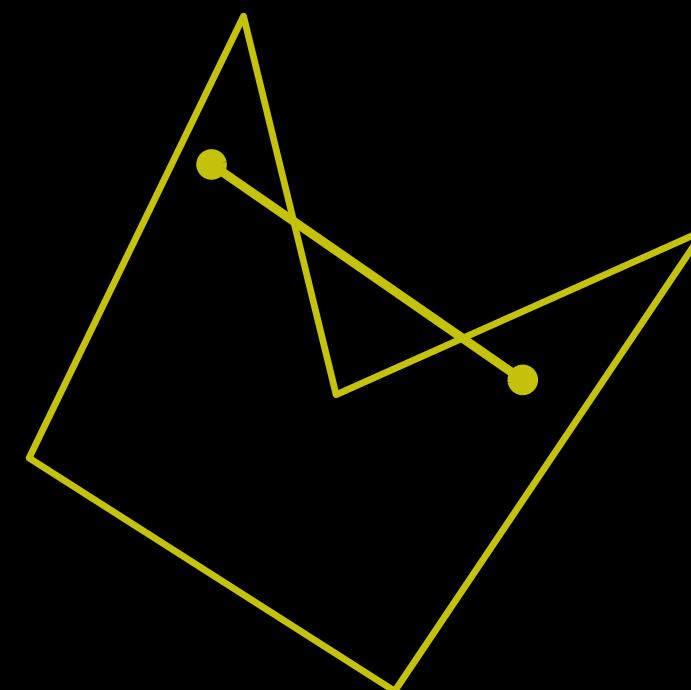
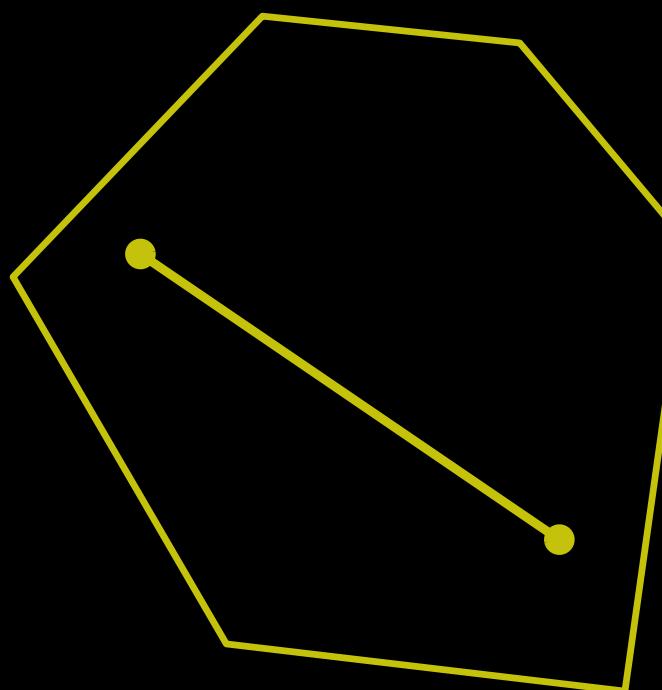
navigation meshes

A* at Scale

- As your space gets bigger and the obstacles more complex, A* (even though it is great) can take too long
 - Even worse, if you have lots of entities doing pathfinding
- Navigation meshes (navmesh) comes from robotics research, 1986

Convex Polygons

- A polygon where **any** line segment connecting two points on the interior or boundaries ...
- ... contains only points on the interior or boundaries



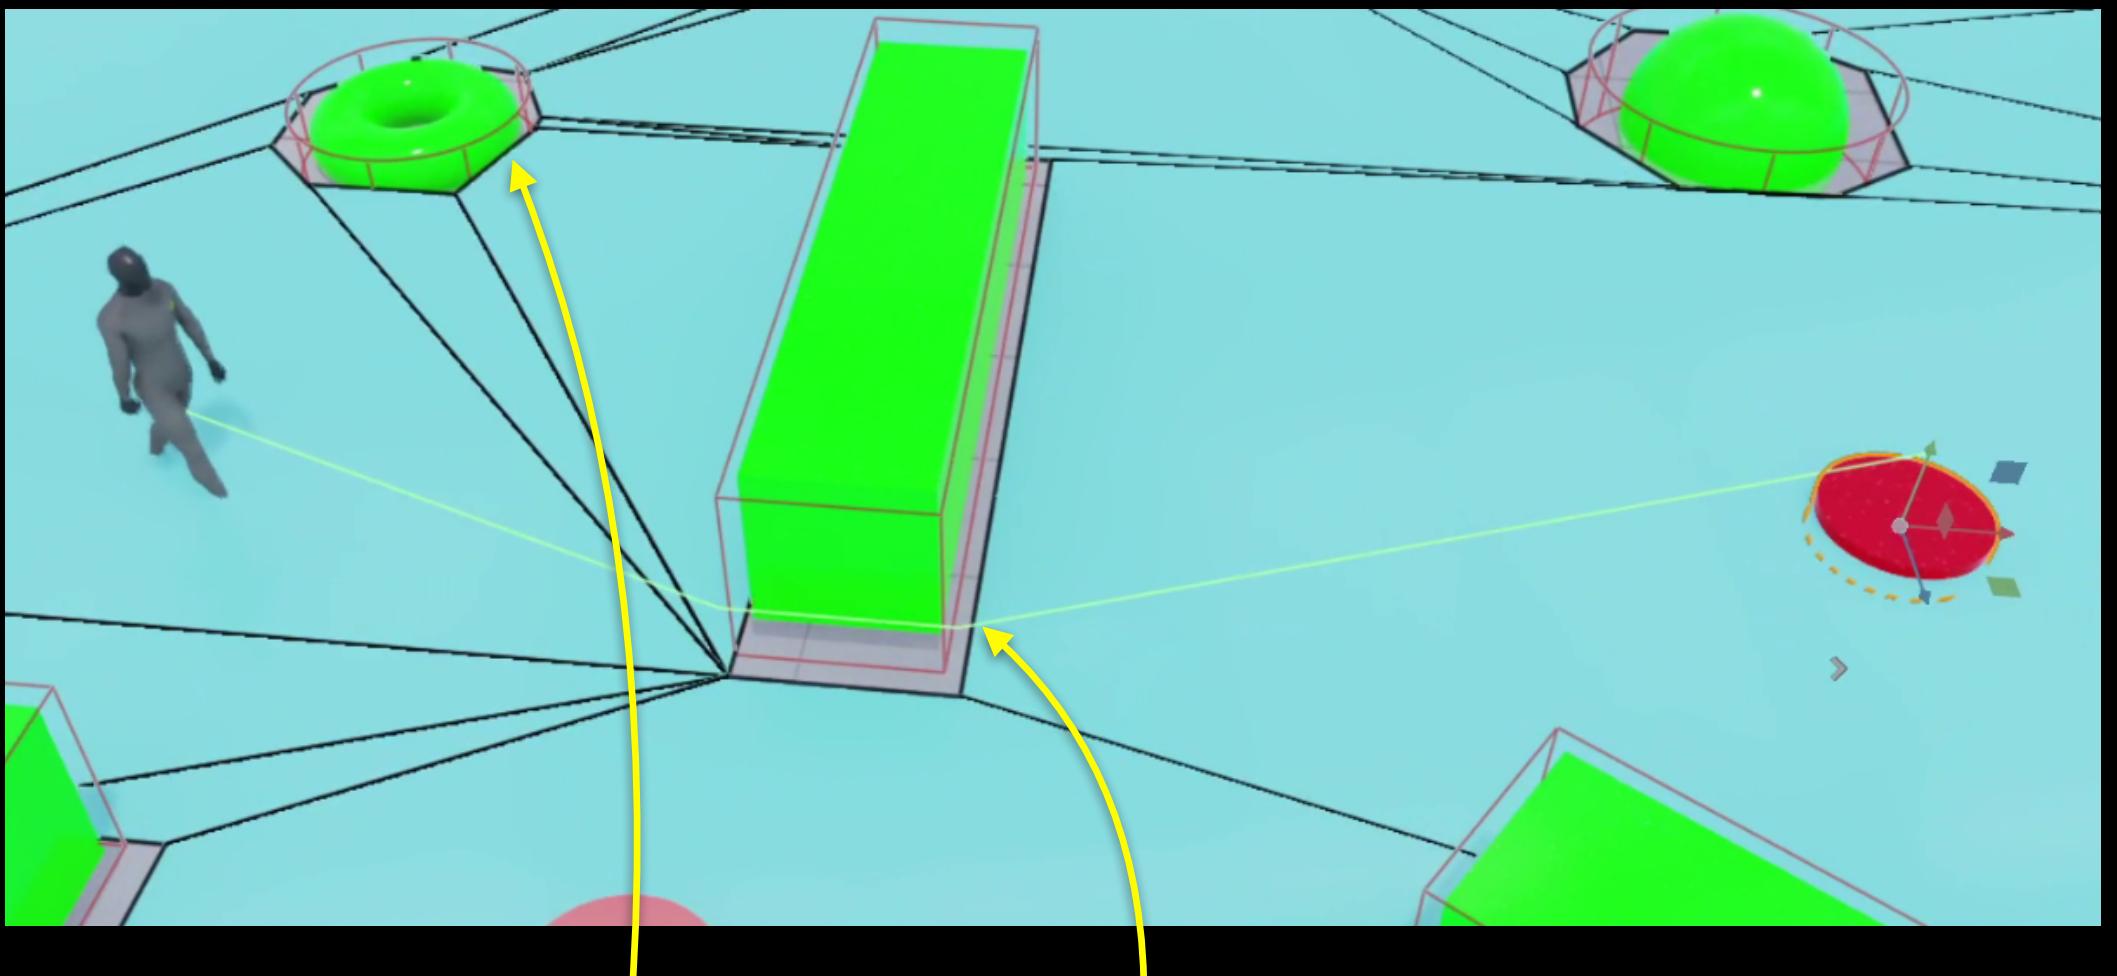
Properties

- Convex polygons have lots of interesting properties that make them useful
- Easy to navigate between any two points in the polygon
 - Just go in a straight line
 - Guarantees you will remain within the polygon

Navmesh

- A collection of convex polygons that define all freely-traversable areas in your game
 - i.e. every space where an agent (player, monster, etc) can legally walk
- All obstacles exist outside of the navmesh
- All polygons are connected by sides where traversal is allowed

Example



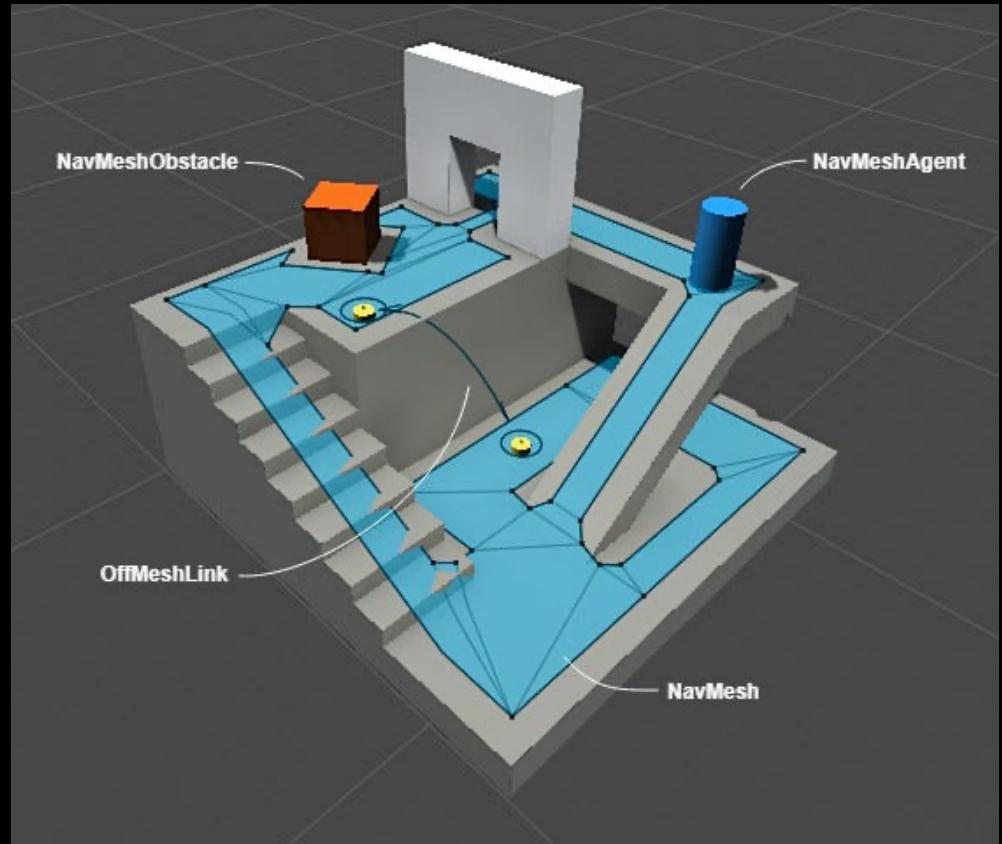
Notice that obstacles are not included -- they are just a hole in the mesh

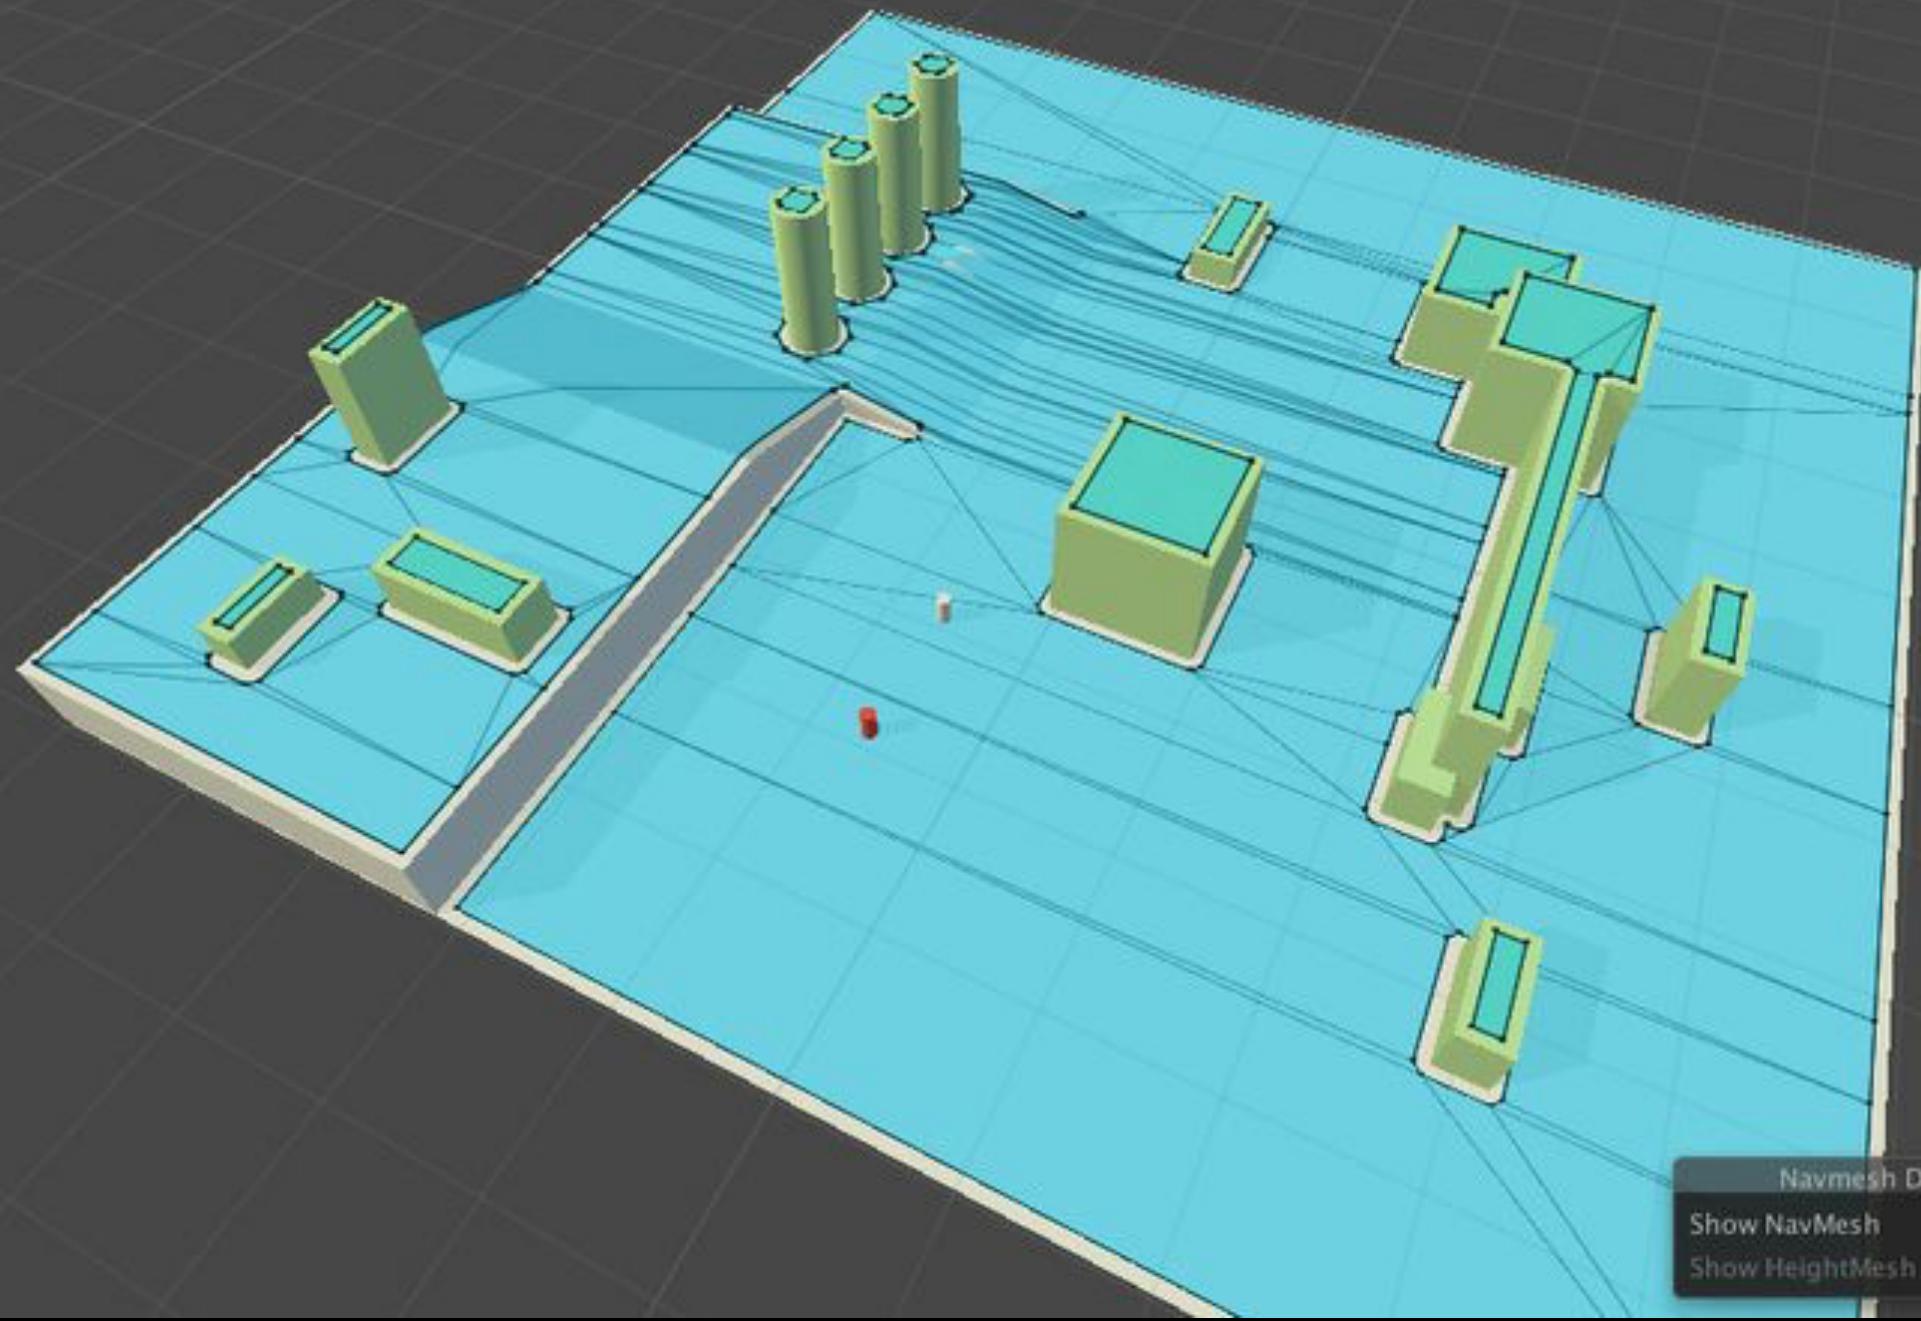
Navigation in the Mesh

1. If the target is in the same polygon as the agent
 - Just go in a straight line
2. If the target is not in the same mesh
 - The connection of polygons form a graph
 - Use A* to determine which polygons to travel through to get to the target's polygon
 - Then use rule #1

Added Benefits

- Each polygon must be 2-dimensional
- But, the collection of polygons doesn't have to be
 - Exploit to model platforms, mezzanines, stairs, different levels, etc





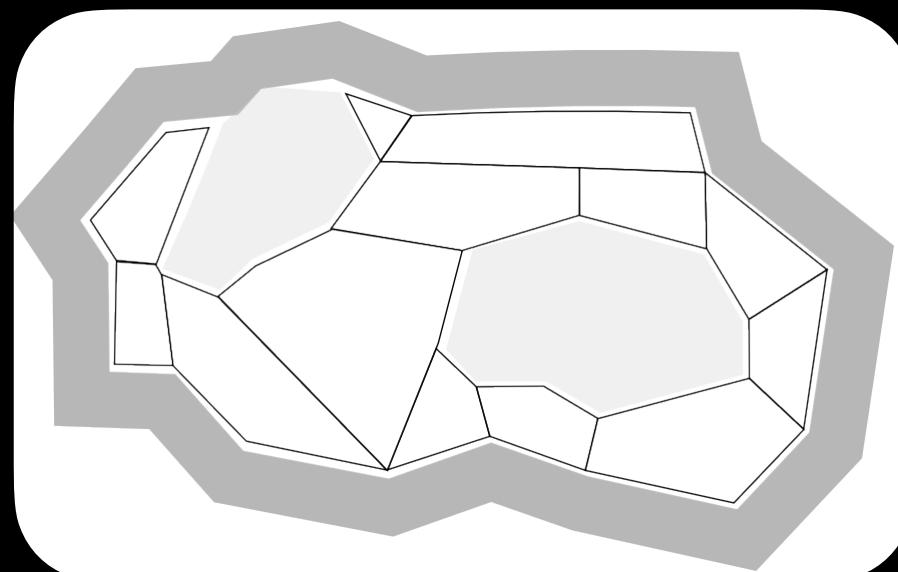
funnel algorithm

Funnel Algorithm

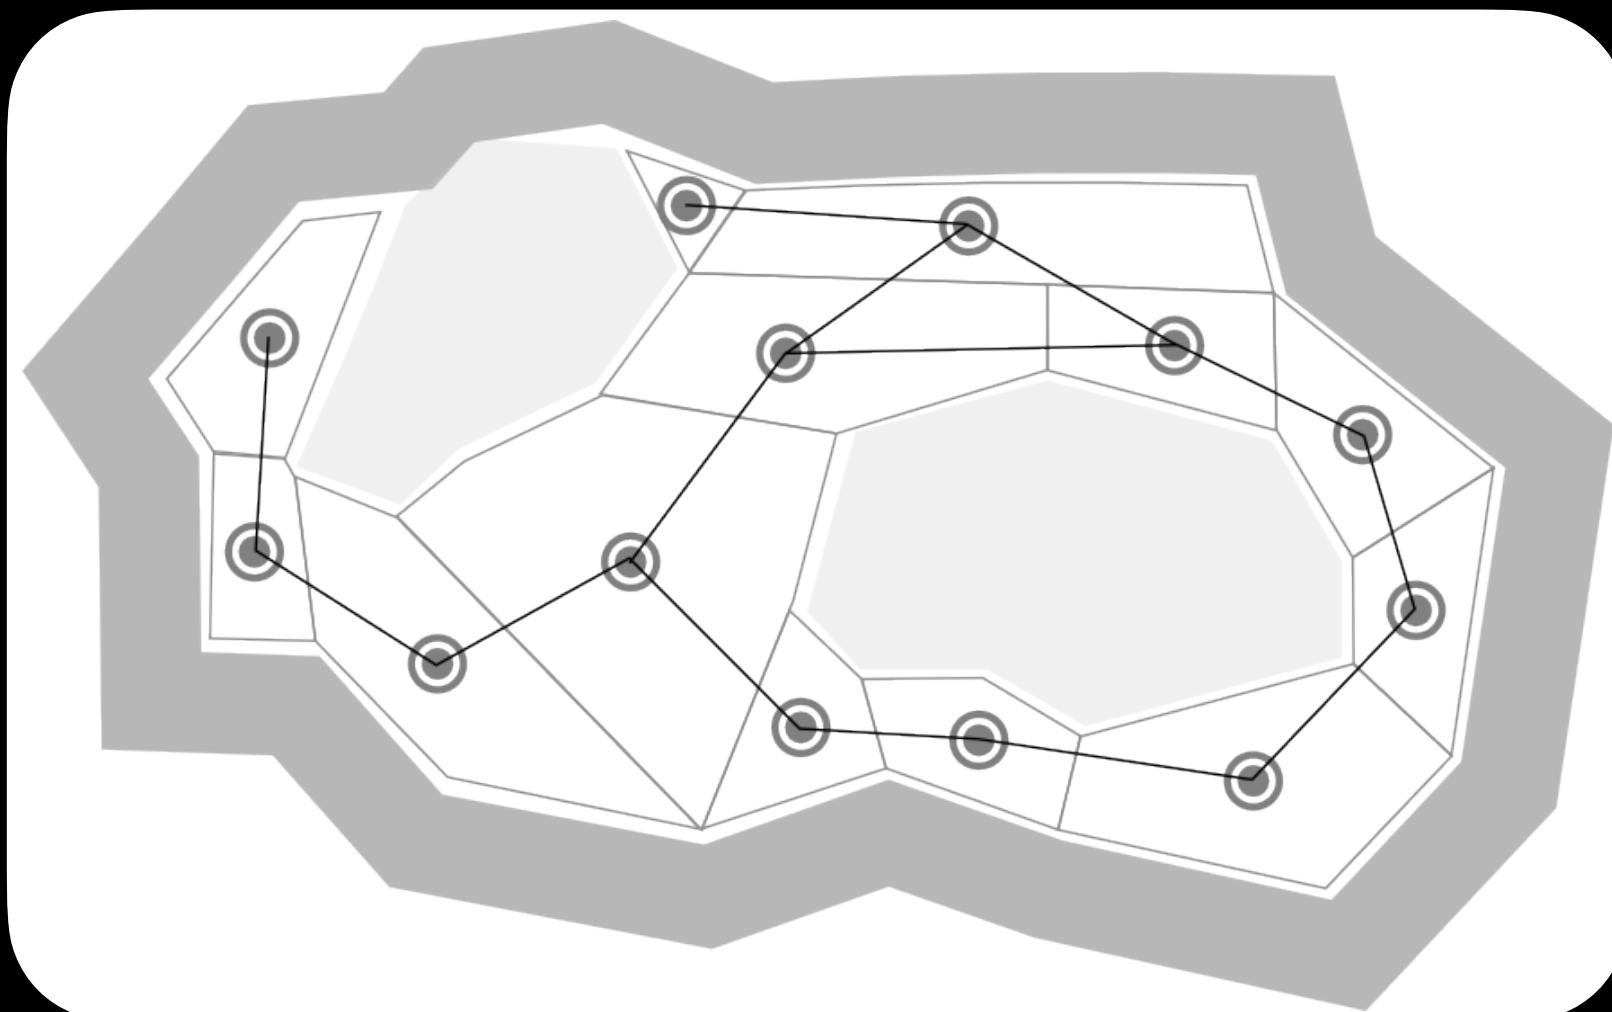
- The navmesh is an abstraction of the complete geography of your game
- Navigation calculations are simplified
- ... but, the resulting path is also simplified and may not be realistic
- Funnel algorithm: improving the resulting path

Funnel Algorithm in Pictures

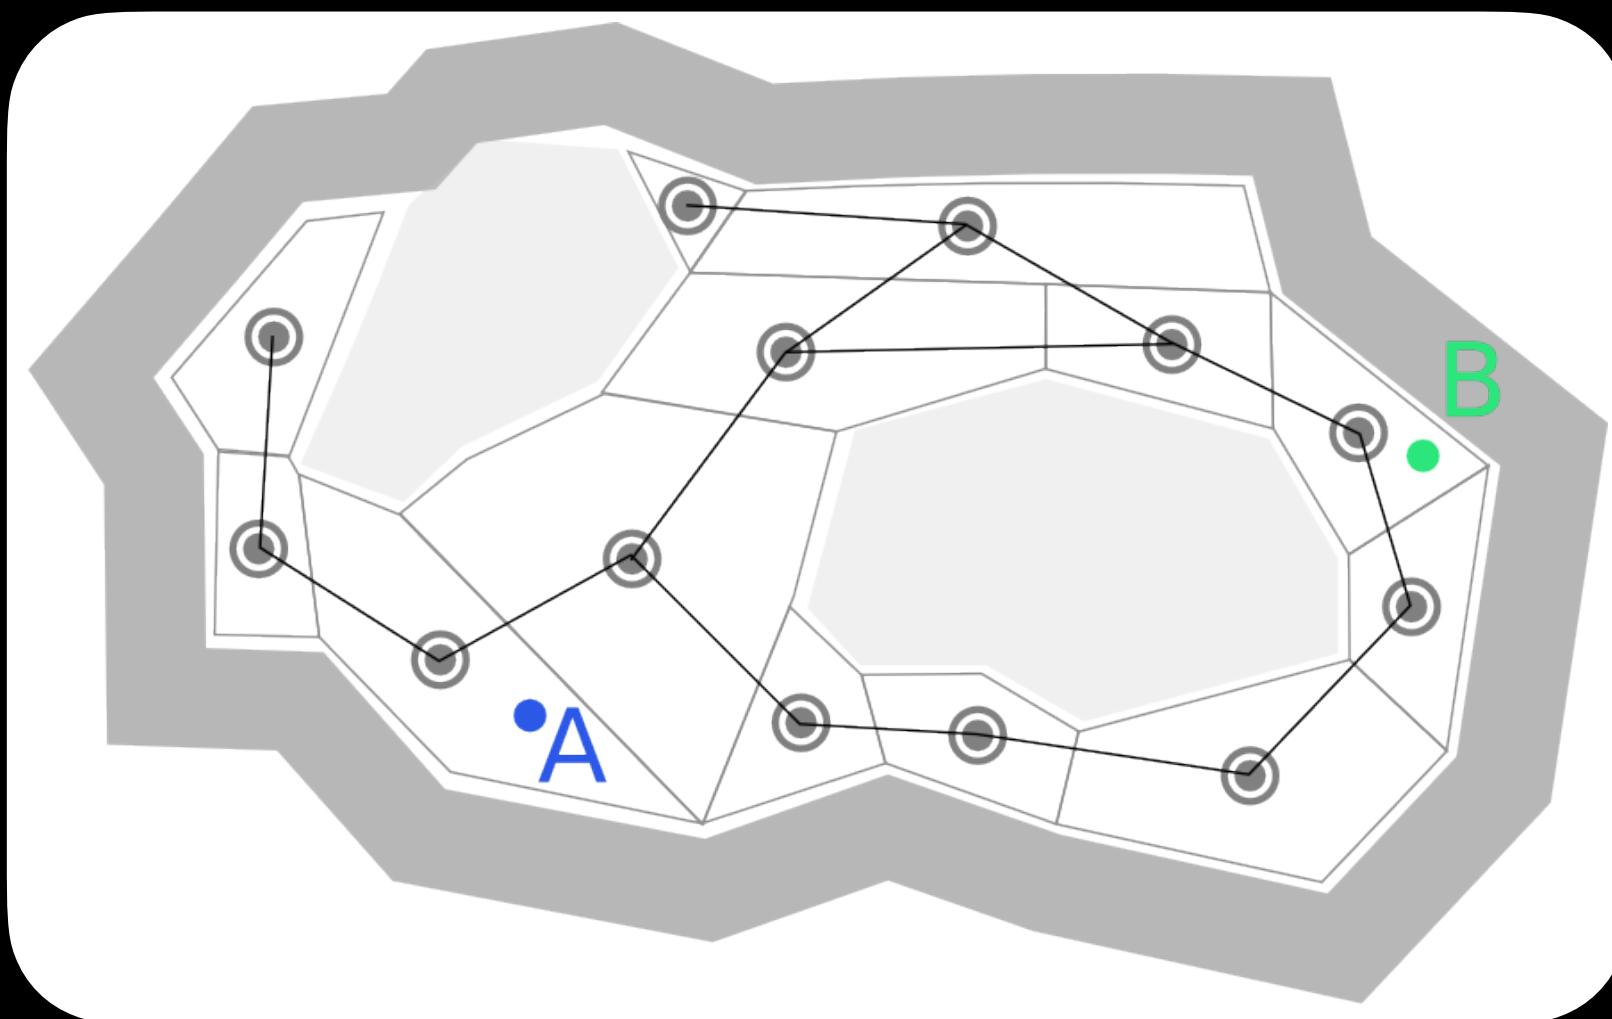
- I took this example from jceipek.com/Olin-Coding-Tutorials/pathing.html
- Imagine an island with two hills (prohibited movement) and the navigable land broken into 13 convex polygons



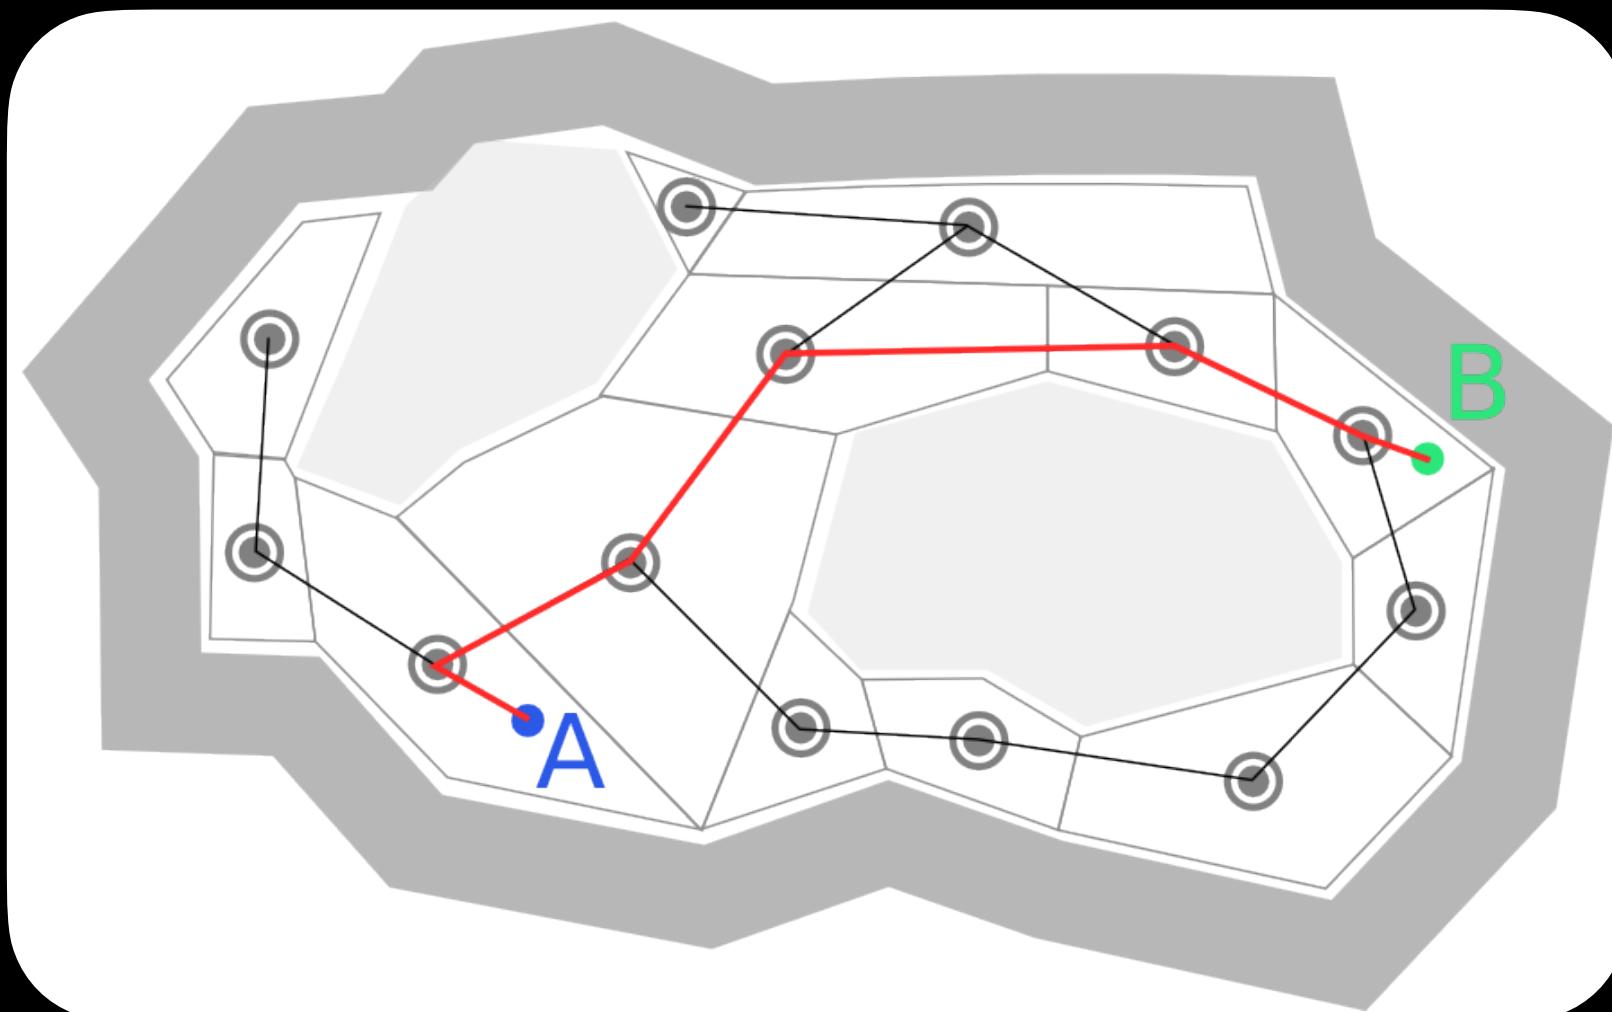
- The navmesh is a graph, so let's put a vertex in each polygon and draw edges between adjacent polygons



- Travel between A and B involves multiple convex polygons
- Must use Rule #2 for navigation (not straight line)



- A* tells us to go along the path of these polygons
- But, it isn't smoothed



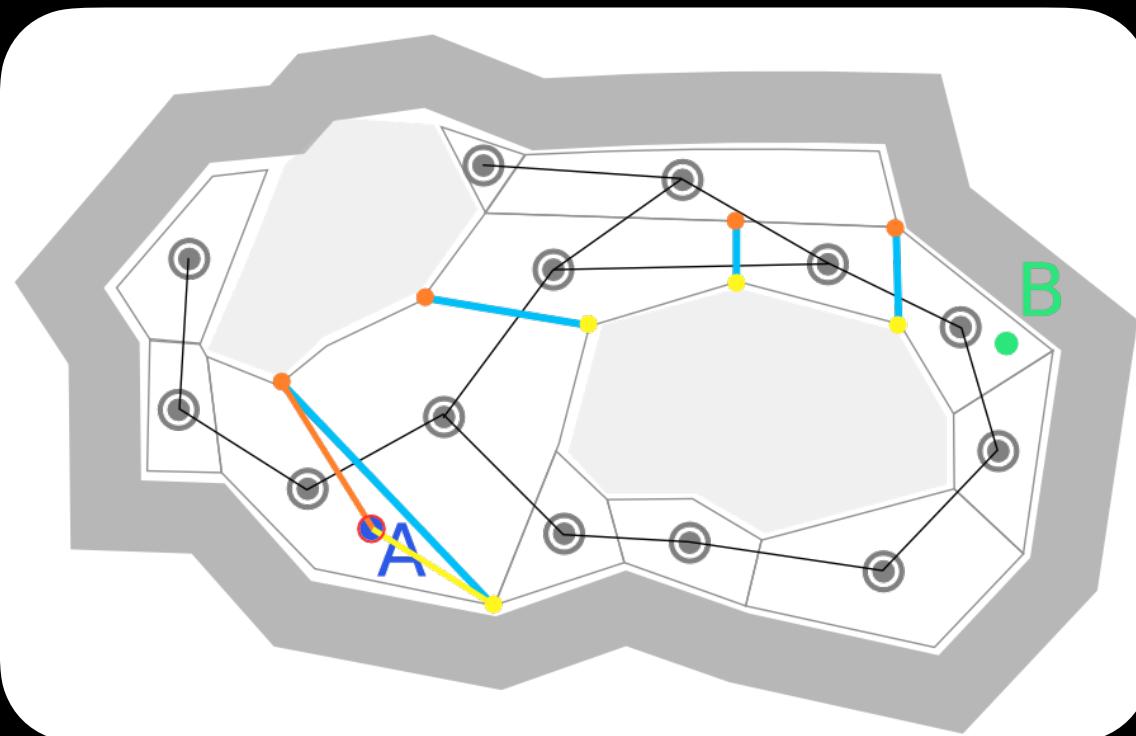
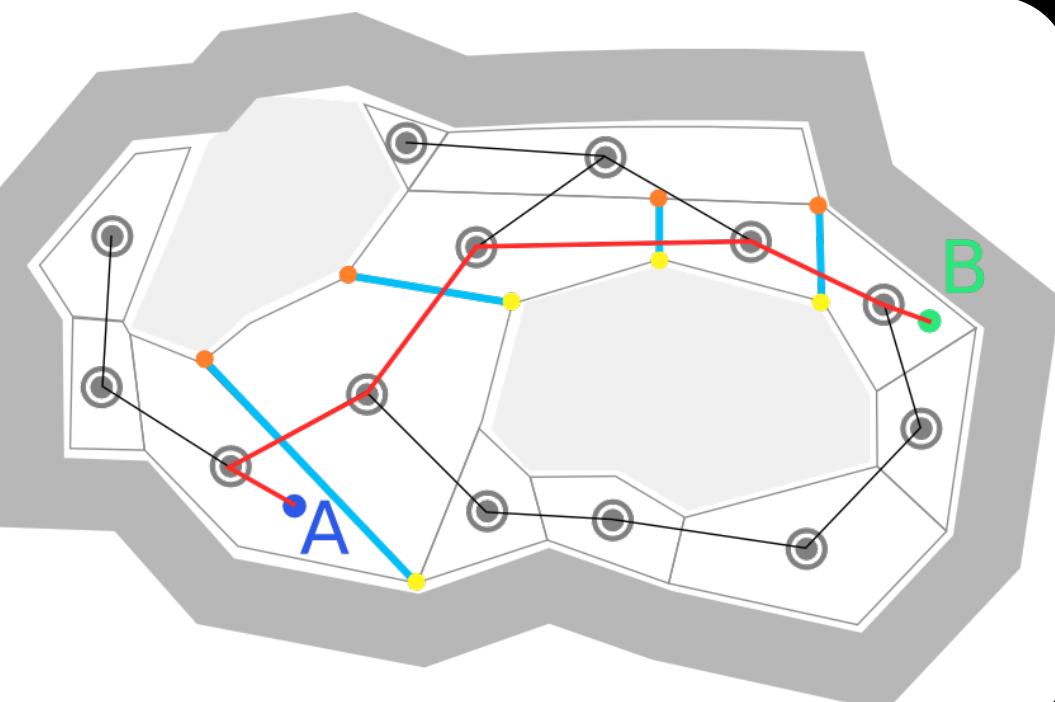
The Algorithm

- Start at A (not the vertex in A's polygon)
- Consider the two edges of the common line segment between polygons (called a "portal")
 - Remember, A* told you which is the next polygon to move into
- Create a "funnel" with three points
 - The apex (point A), left portal end, right portal end

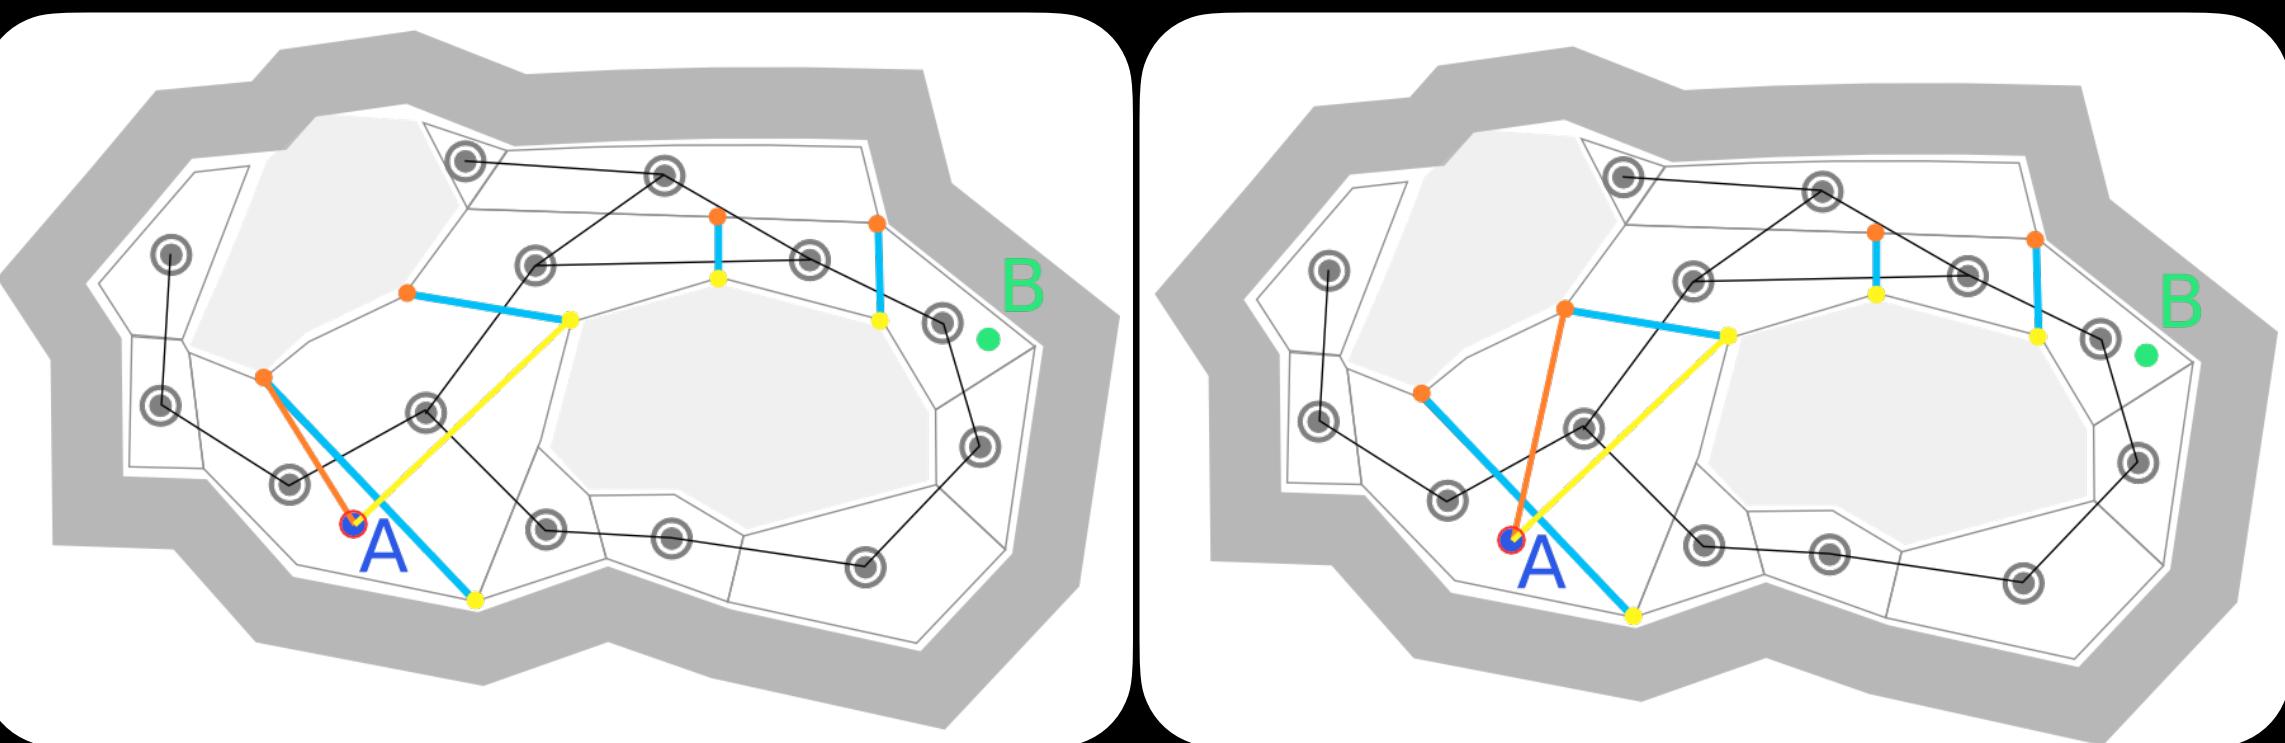
- Now, consider the second portal
- Move the left point of the funnel to the left edge of this second portal
 - If the update would widen the funnel, then ignore the update
 - Remember: you're still looking for a straight line to travel from the apex, so you can't widen or you might need to go through a corner
- Repeat with the right of the funnel

- If the update would crosses the sides of the funnel (left-over-right or right-over-left)
 - Move the apex to the last point on that side of the funnel and start over
 - Keep track of where the apex has been, as that will be your path

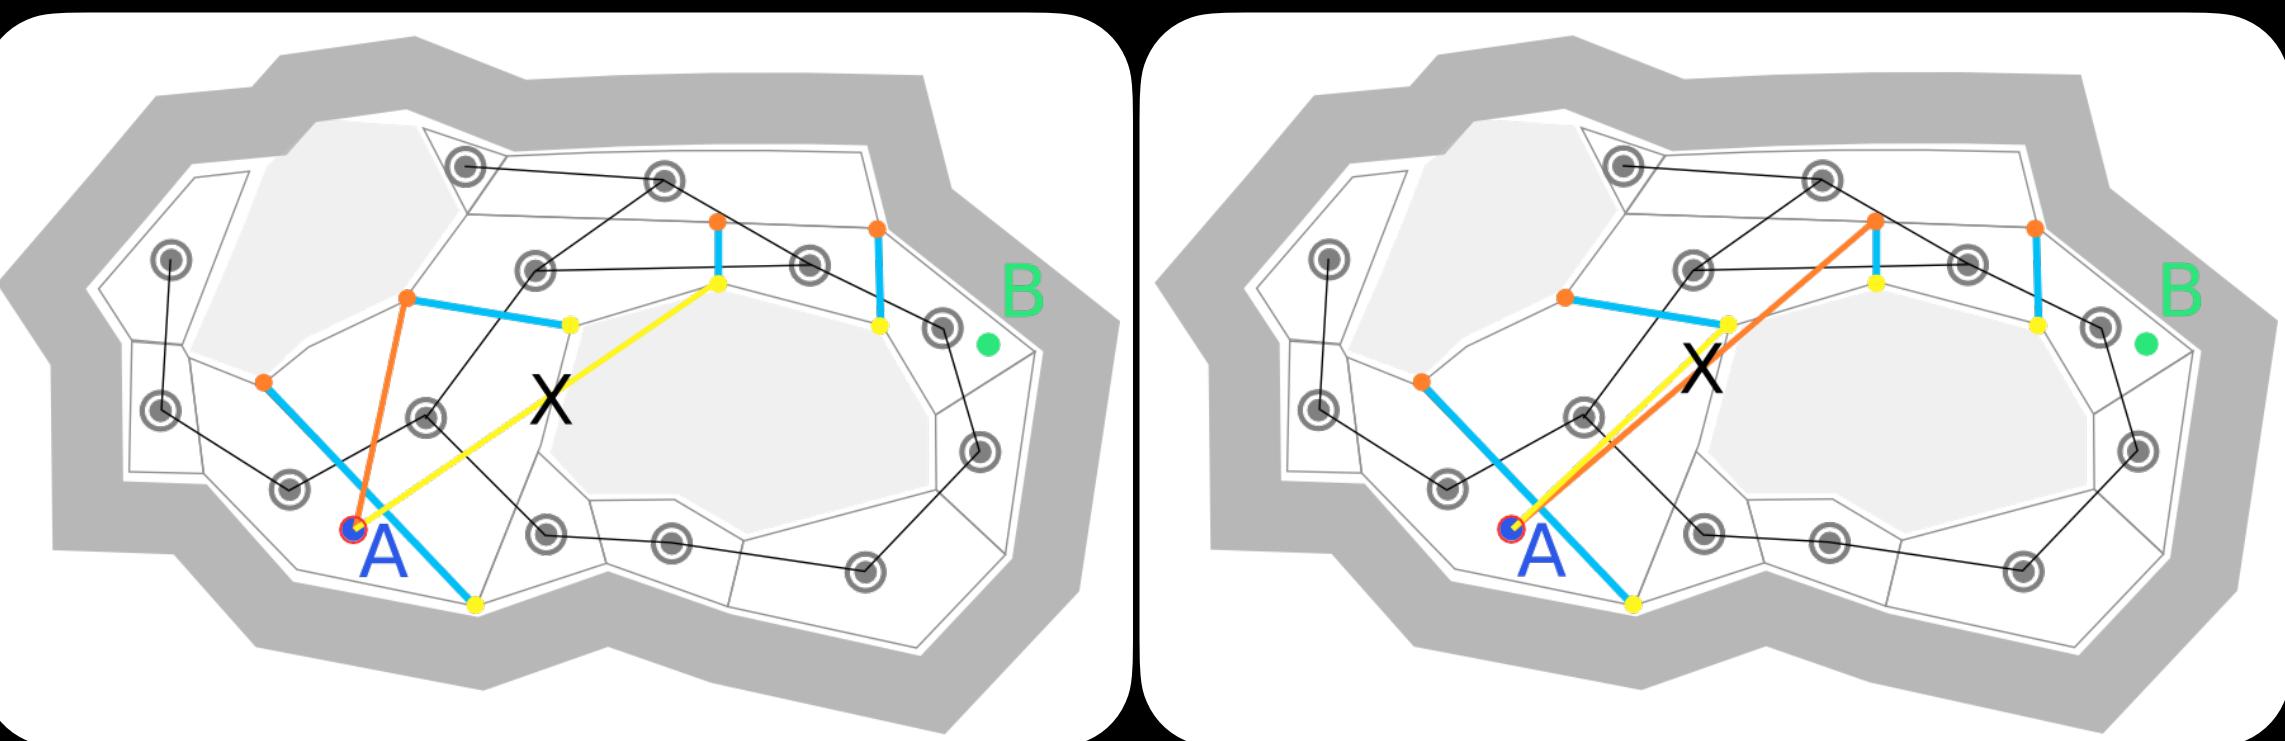
Pictures



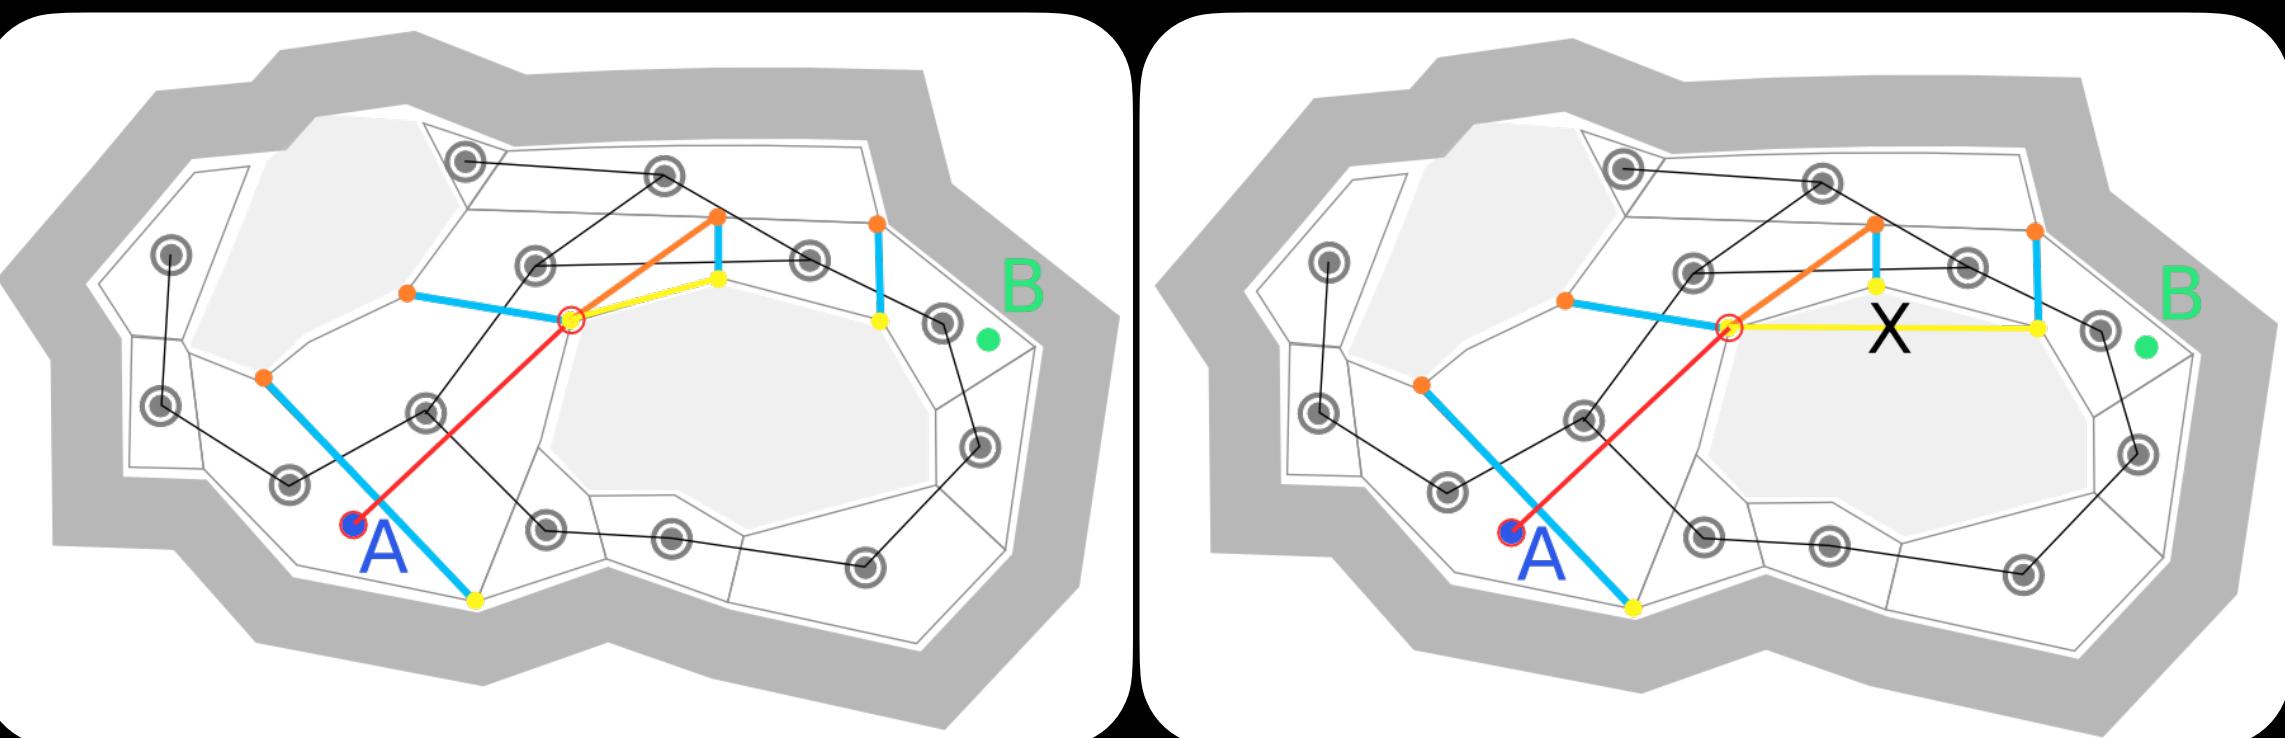
Pictures



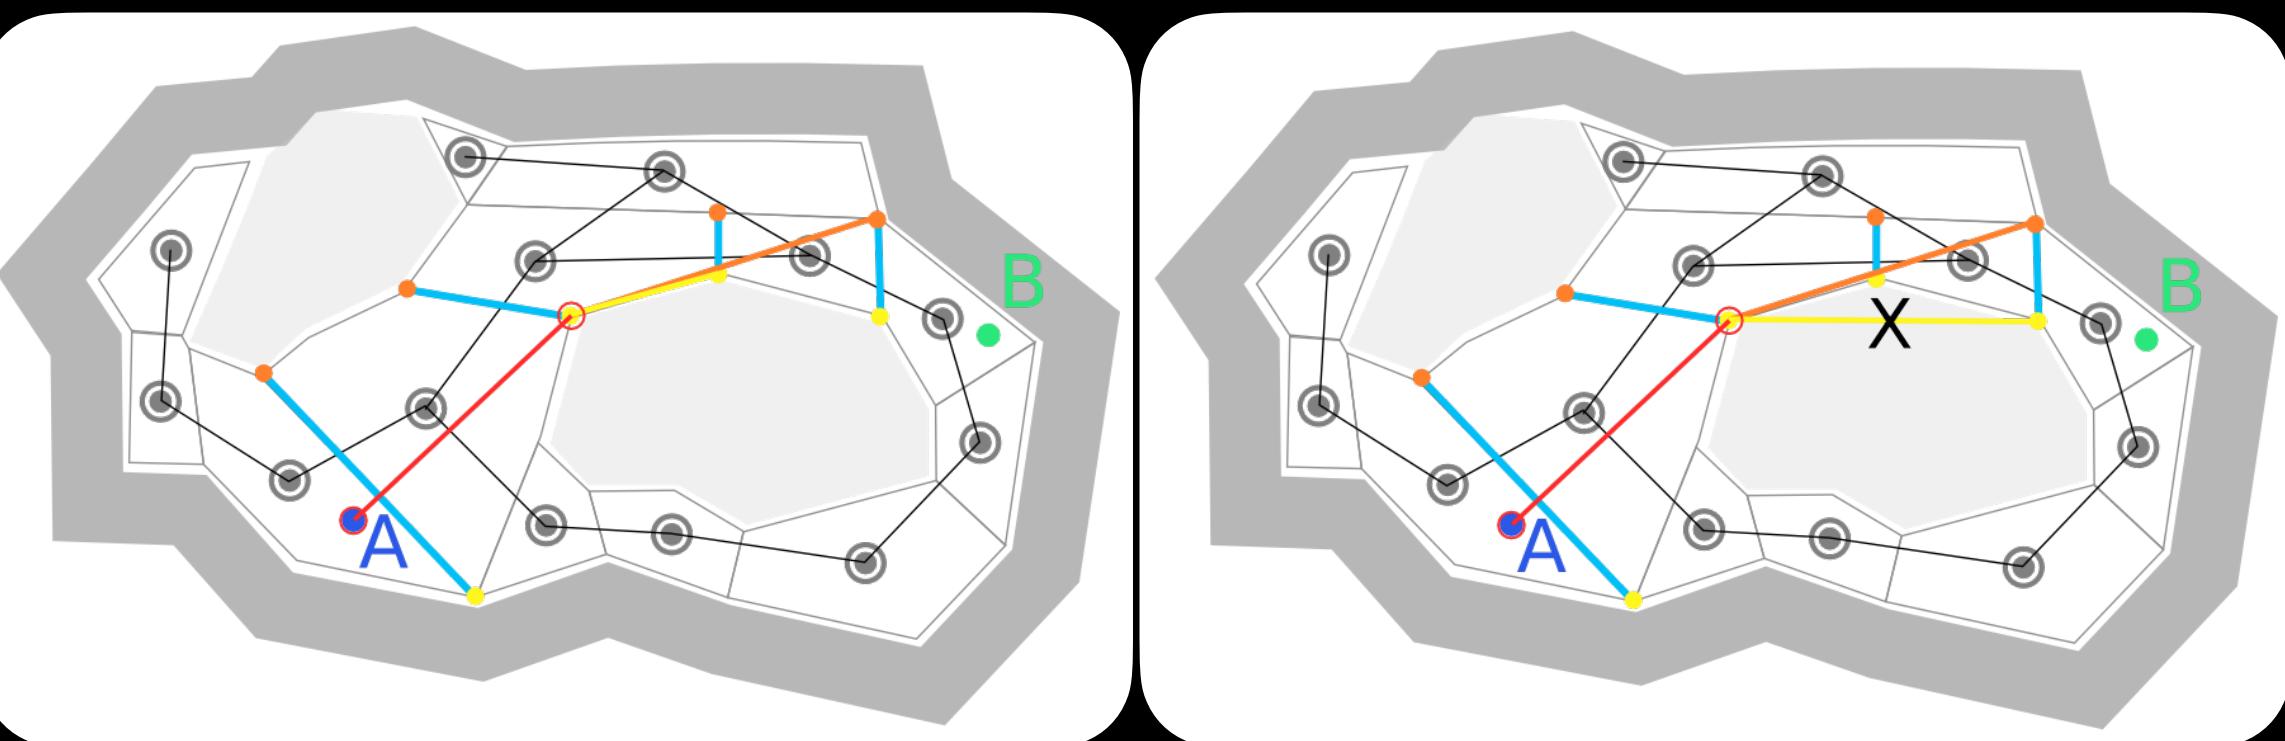
Pictures



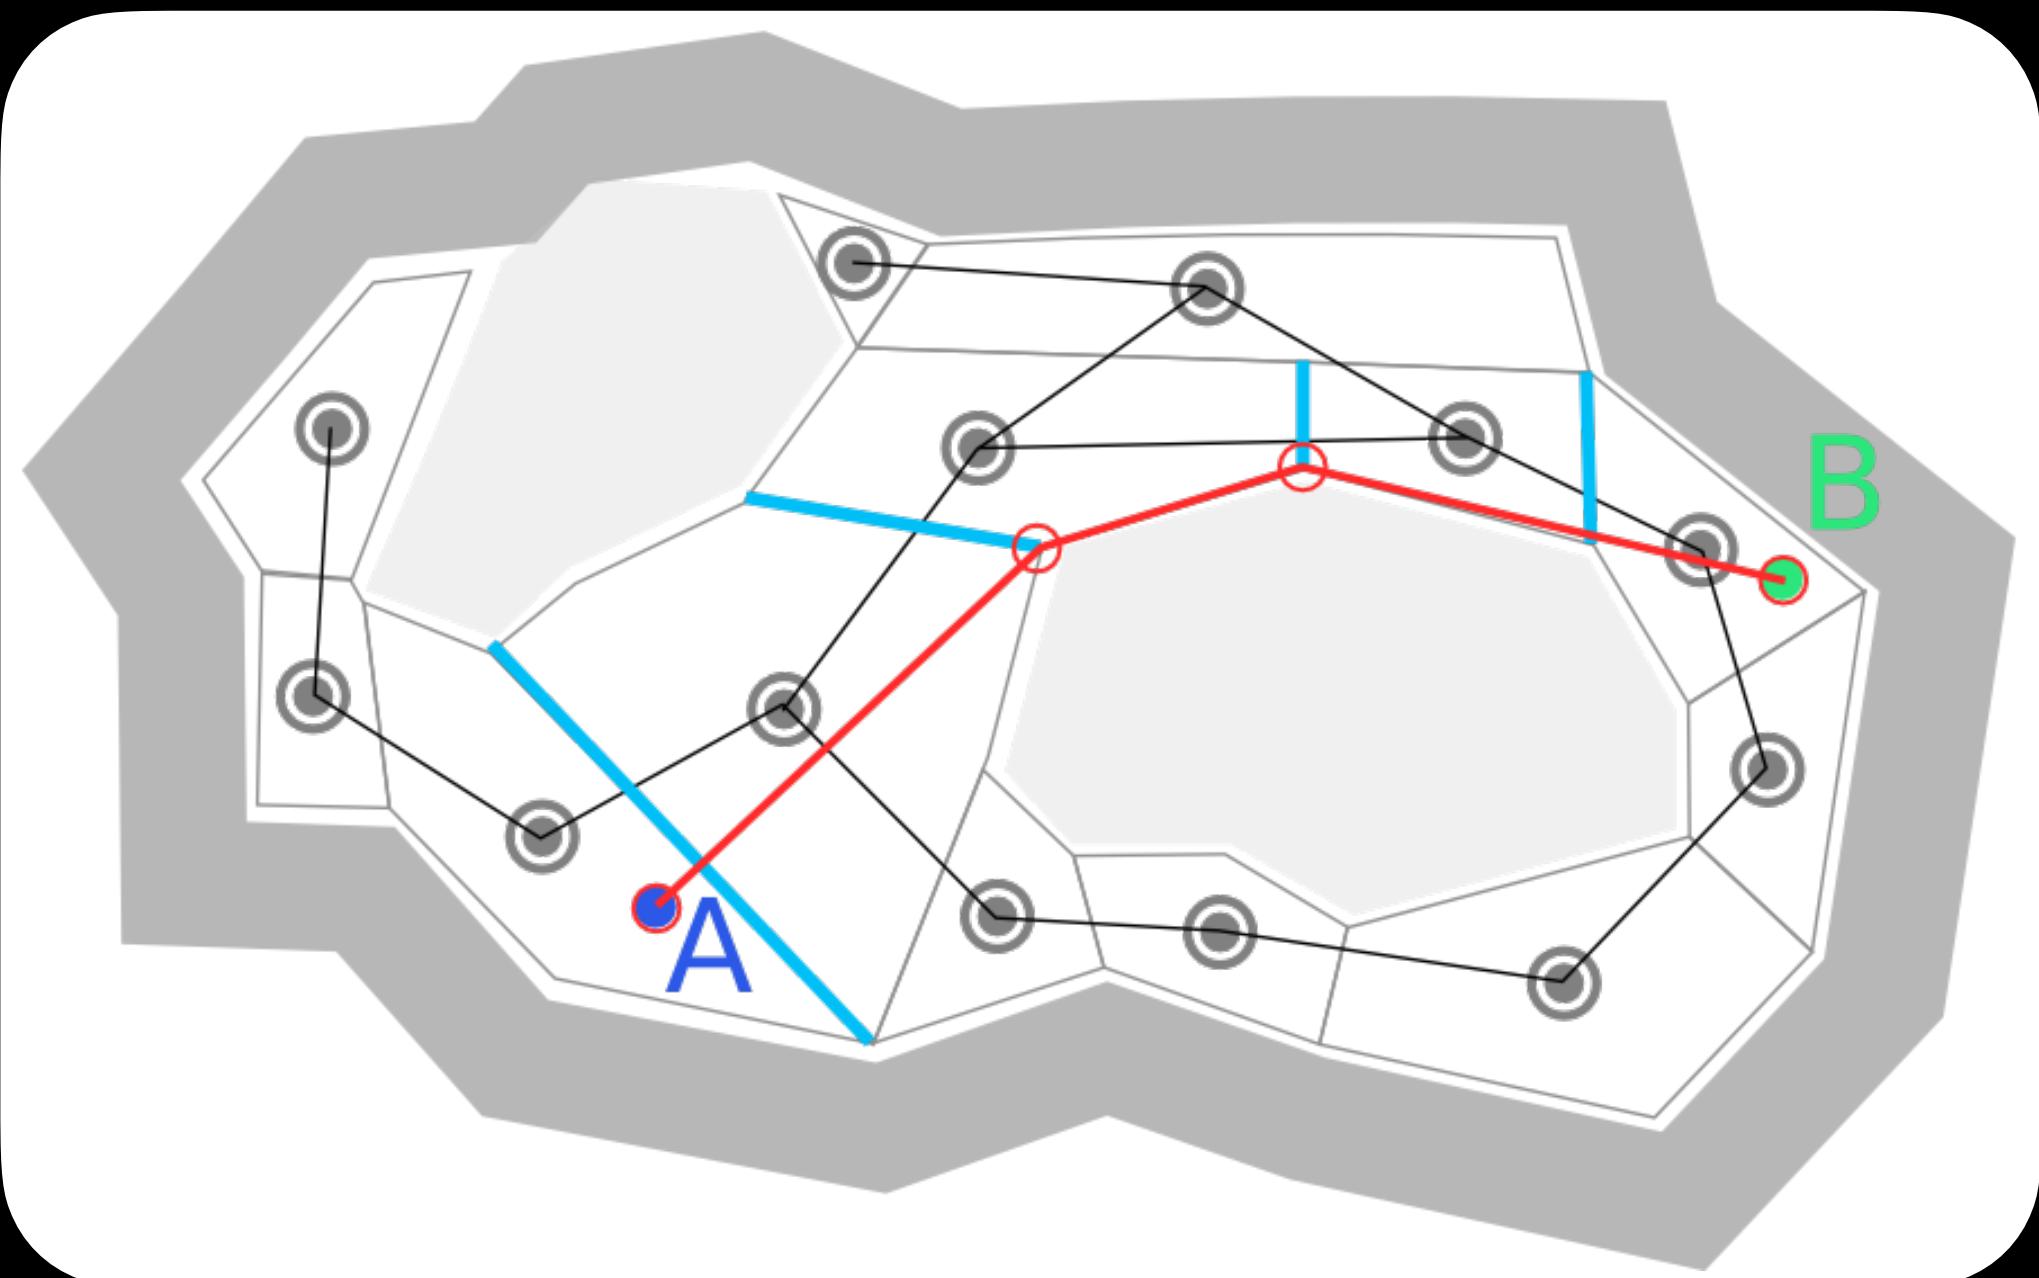
Pictures



Pictures



Pictures



What did you learn today?

- Breadth First Search: Flooding the entire grid
- Pathfinding: Keeping track of where you came from
- Early exit: If you have a target, you can stop when you get to it.
- Dijkstra's Algorithm (again): Least cost path from source to every cell. But inefficient
- Greedy Best First Search (Heuristic): Very efficient, except if obstacles are in the way
- A* Search (Heuristic): a mix of Dijkstra's with a heuristic. Best of both worlds
- Navigation Meshes: When A* would be too expensive to run, we can abstract it.
- There are a lot of Navigation algorithms (and there are many more than we've discussed).